COMP 3234B
Computer and Communication Networks

2nd semester 2023-2024

Transport Layer (II)

Prof. C. Wu

Department of Computer Science
The University of Hong Kong

Roadmap

Transport layer

- Principles behind transport-layer services
  - multiplexing/demultiplexing (ILO1, 2)
  - reliable data transfer:(ILO 2, 3)
    - rdt 1.0
    - rdt 2.0, 2.1, 2.2
    - rdt 3.0
    - GBN
    - selective repeat
  - flow control (ILO 2, 3)
  - congestion control (ILO 2, 3)
- Transport protocols in the Internet (ILO 2, 3)
  - TCP
  - UDP

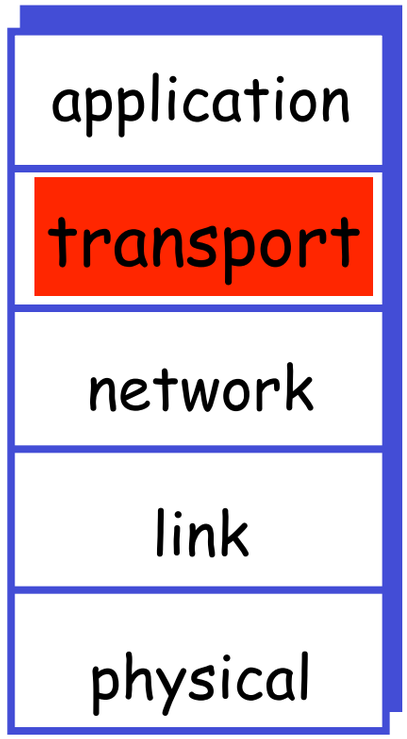

Pipelining

- Sender can send multiple packets without waiting for acknowledgements

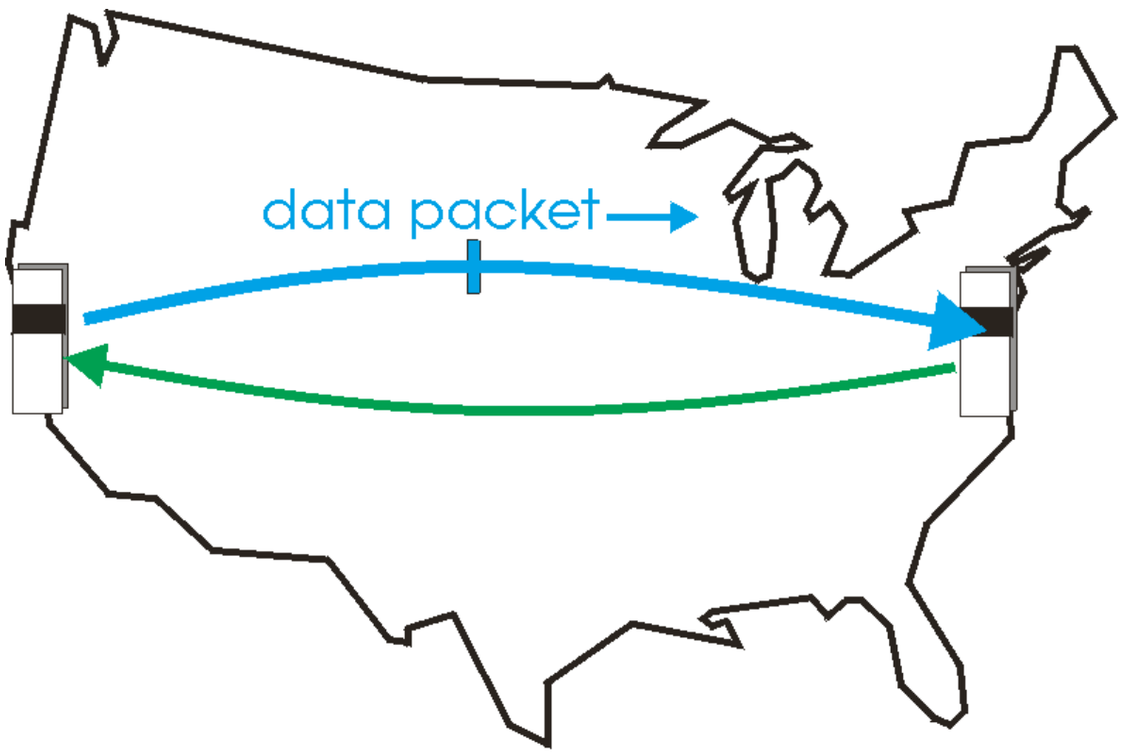

(a) a stop-and-wait protocol in operation

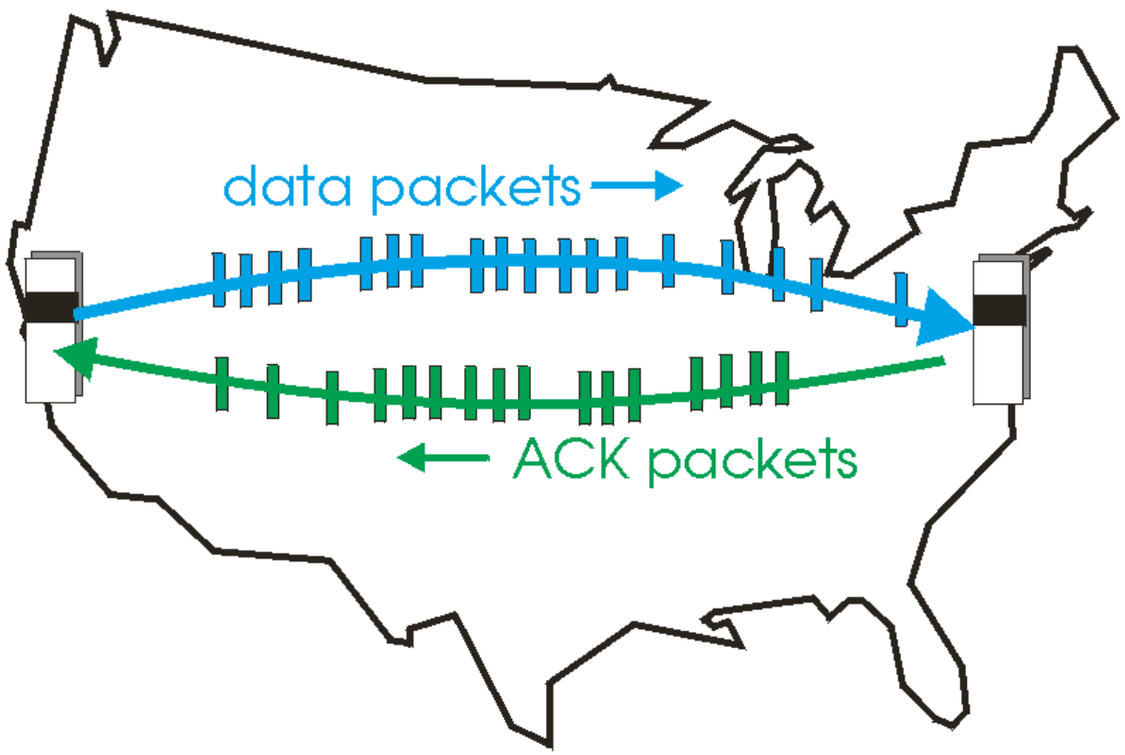

(b) a pipelined protocol in operation

Pipelining: performance

increased utilization

link rate R=1 Gbps, propagation delay 15 ms, each packet L= 8000 bit

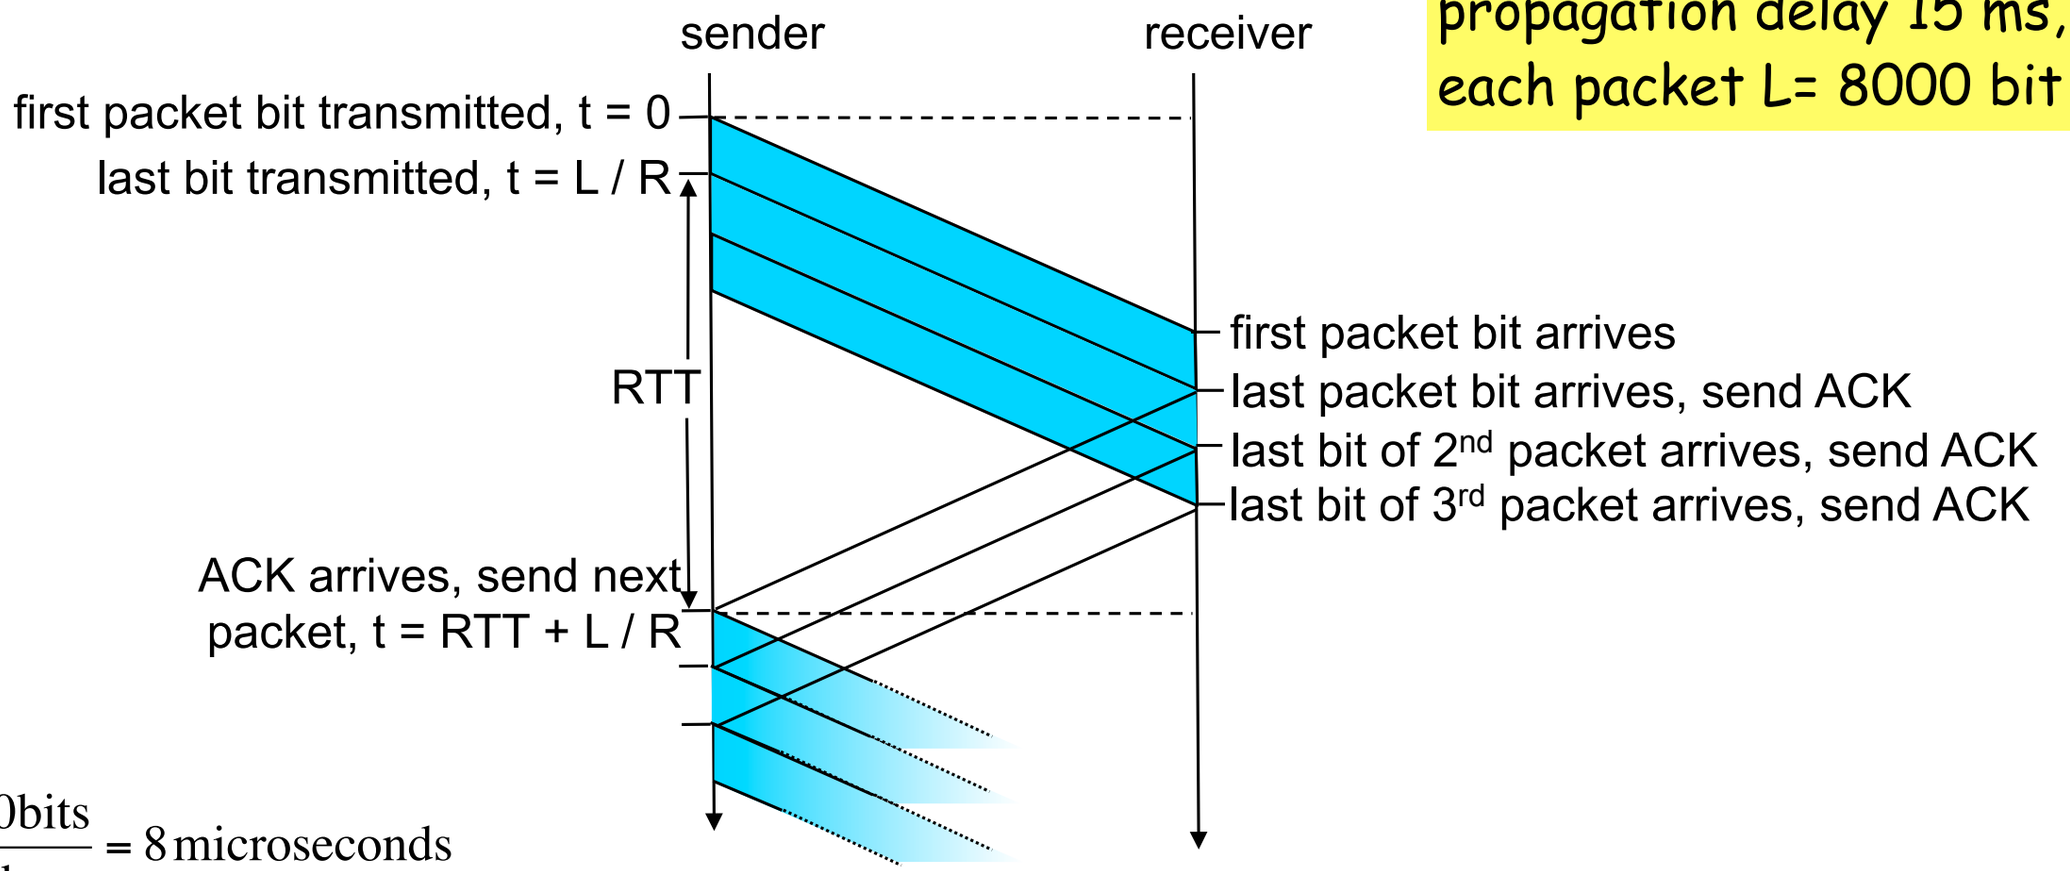

$$d_{trans} = \frac{L}{R} = \frac{8000\text{bits}}{10^9\text{bps}} = 8\,\text{microseconds}$$

$$RTT = 2 \times 15ms = 30ms$$

3-packet pipelining increases utilization by a factor of 3!

$$U_{sender} = \frac{3 * L / R}{RTT + L / R} = \frac{.024}{30.008} = 0.0008$$

Pipelining protocols

- New requirements for rdt
  - range of sequence numbers must be increased
  - buffering more than one packet at sender and/or receiver

- Two generic forms of pipelined reliable data transfer protocols

Go-back-N

- Sender can have up to N unacked packets in pipeline
- Receiver only sends cumulative acks
- Sender has timer for oldest unacked packet
  - If timer expires, retransmit all unacked packets

Selective Repeat

- Sender can have up to N unacked packets in pipeline
- Receiver acks individual packets

- Sender maintains timer for each unacked packet
  - When timer expires, retransmit only unack packet

Go-Back-N (GBN): sender

- Sender protocol
  - *k*-bit seq # carried in a field in packet header

    seq # range: 0 to $2^k - 1$
  - sender allows up to *N* consecutive unack'ed packets
  - sender has timer for the oldest unacked packet

    `timeout(n)`: retransmit packet `n` and all packets in window with higher seq #

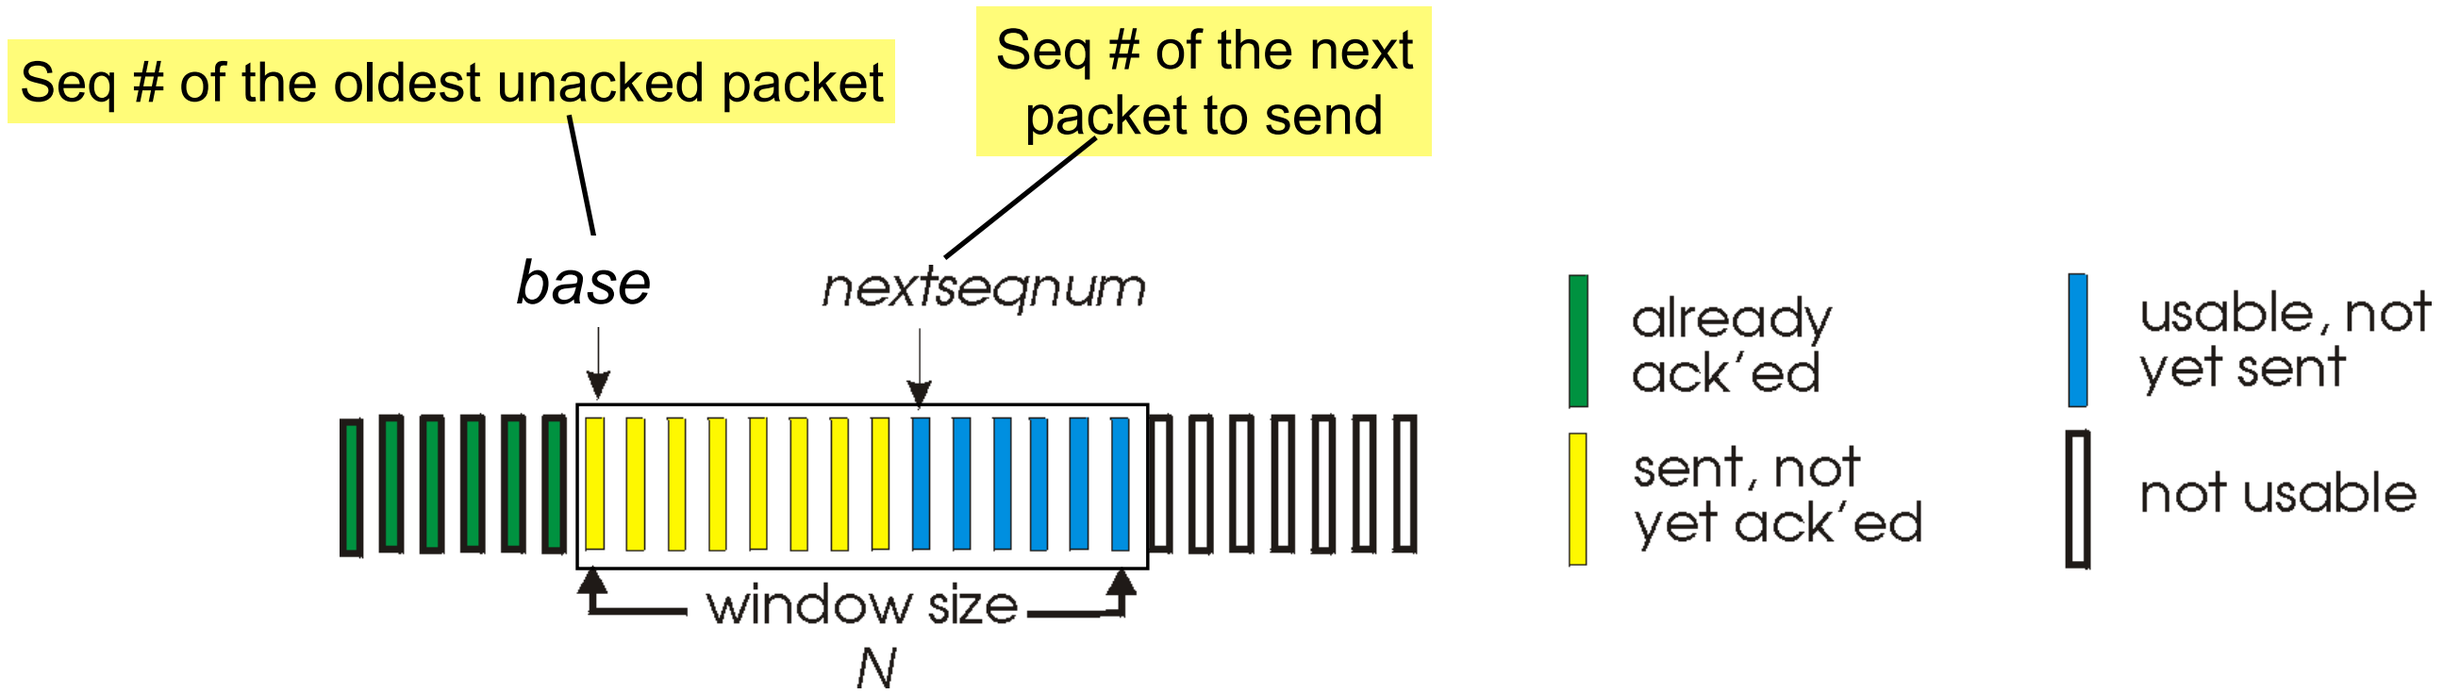

Go-Back-N (GBN): sender (cont'd)

- Sender protocol
  - 4 intervals of seq #

    [*1, base - 1*]: packets already transmitted and acked

    [*base, nextseqnum-1*]: packets sent but not yet acked

    [*nextseqnum, base+N-1*]: packets that can be sent immediately (when data from upper layer arrive)

    [*base+N,* --]: cannot be used until an unacked packet has been acked

  - Window of size *N* slides forward over the seq # space

GBN: sliding-window protocol

Q: why do we limit the number of unacked packets to *N*?
A: flow control, congestion control

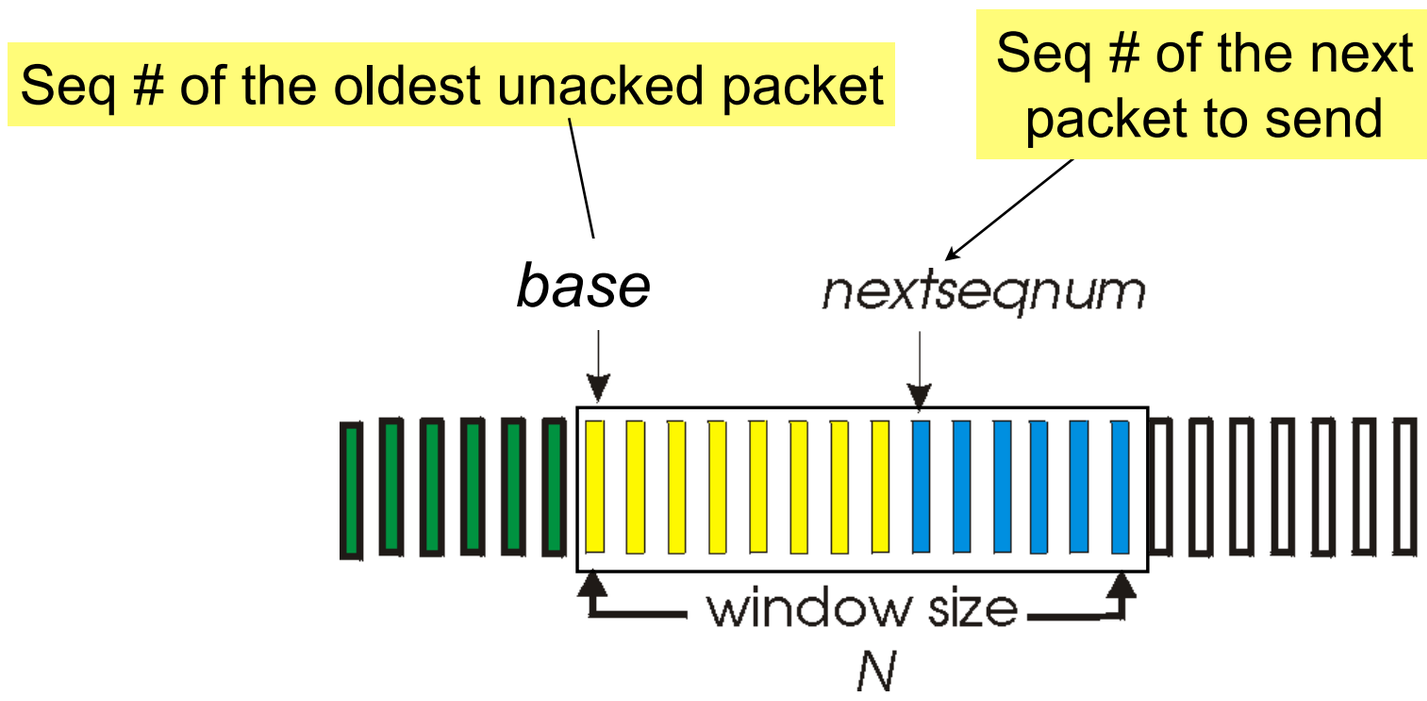

Go-Back-N (GBN): receiver

- Receiver protocol
  - sends ACK for correctly-received packet with highest in-order seq #

    i.e., cumulative ACK

    ACK(n): ACKs all packets up to, including seq # n
  - in any other case, discards packet and resends ACK for the most recently received packet with highest in-order seq #

    e.g., corrupted packet, out-of-order packet

    => may generate duplicate ACKs

Why discarding out-of-order packets? (e.g., if packet *n+1* is received while packet *n* is expected, discard packet *n+1*)

reason: if packet *n* is lost, packet *n+1* will anyway be retransmitted

advantage: no receiver buffering! it need only remember *expectedseqnum*

disadvantage: more retransmissions needed if subsequent retransmission of packet *n+1* is lost or corrupted

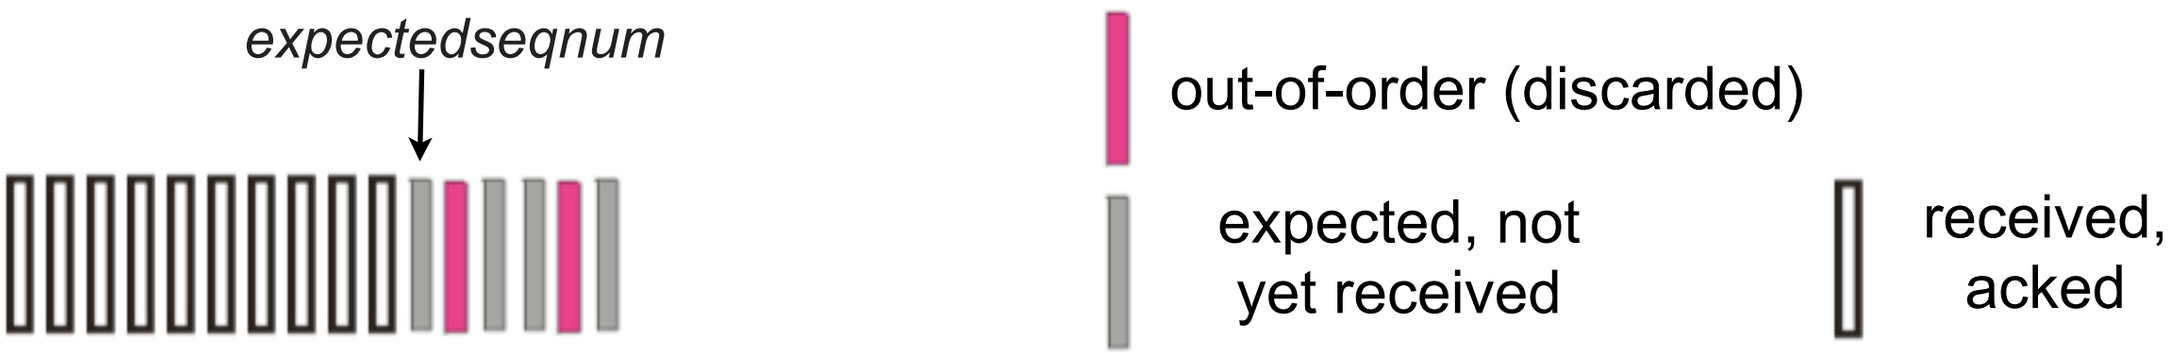

Go-Back-N (GBN): sender FSM

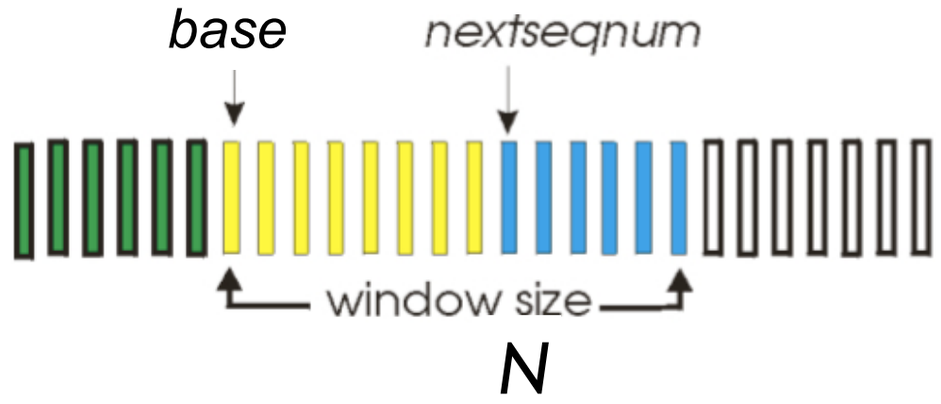

Sender FSM

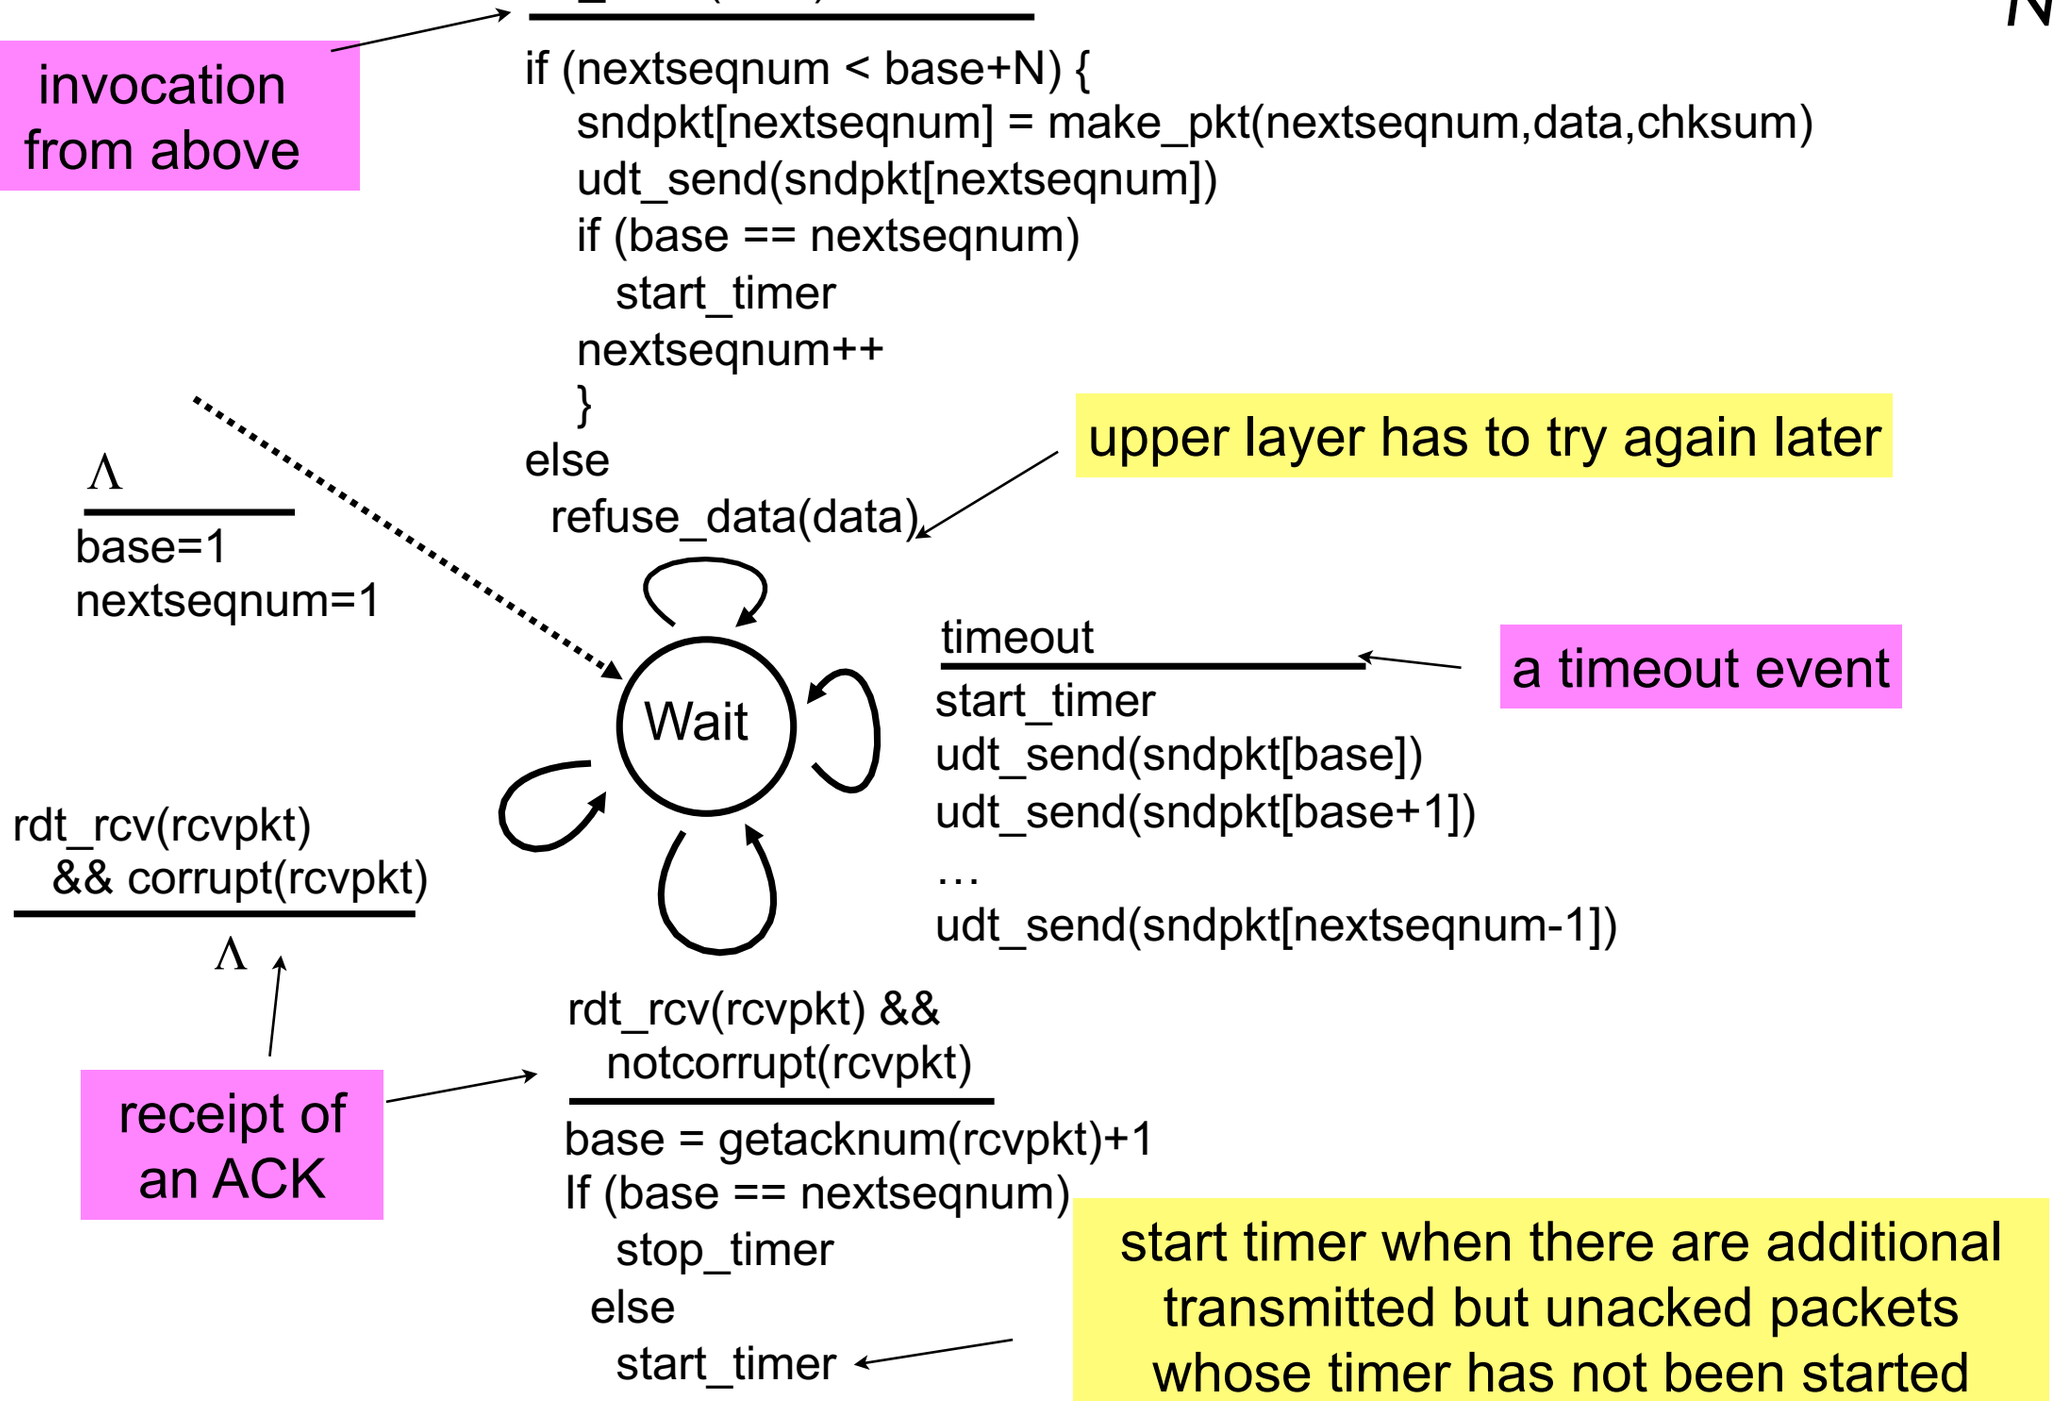

Go-Back-N (GBN): receiver FSM

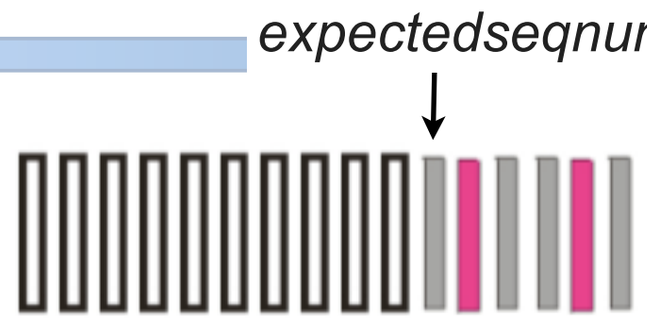

Receiver FSM

any other cases

default

udt_send(sndpkt)

Λ

expectedseqnum=1
sndpkt =
  make_pkt(0,ACK,chksum)

Wait

correctly receiving packet with in-order seq #

rdt_rcv(rcvpkt)
  && notcurrupt(rcvpkt)
  && hasseqnum(rcvpkt,expectedseqnum)

extract(rcvpkt,data)
deliver_data(data)
sndpkt = make_pkt(expectedseqnum,ACK,chksum)
udt_send(sndpkt)
expectedseqnum++

GBN in action (1)

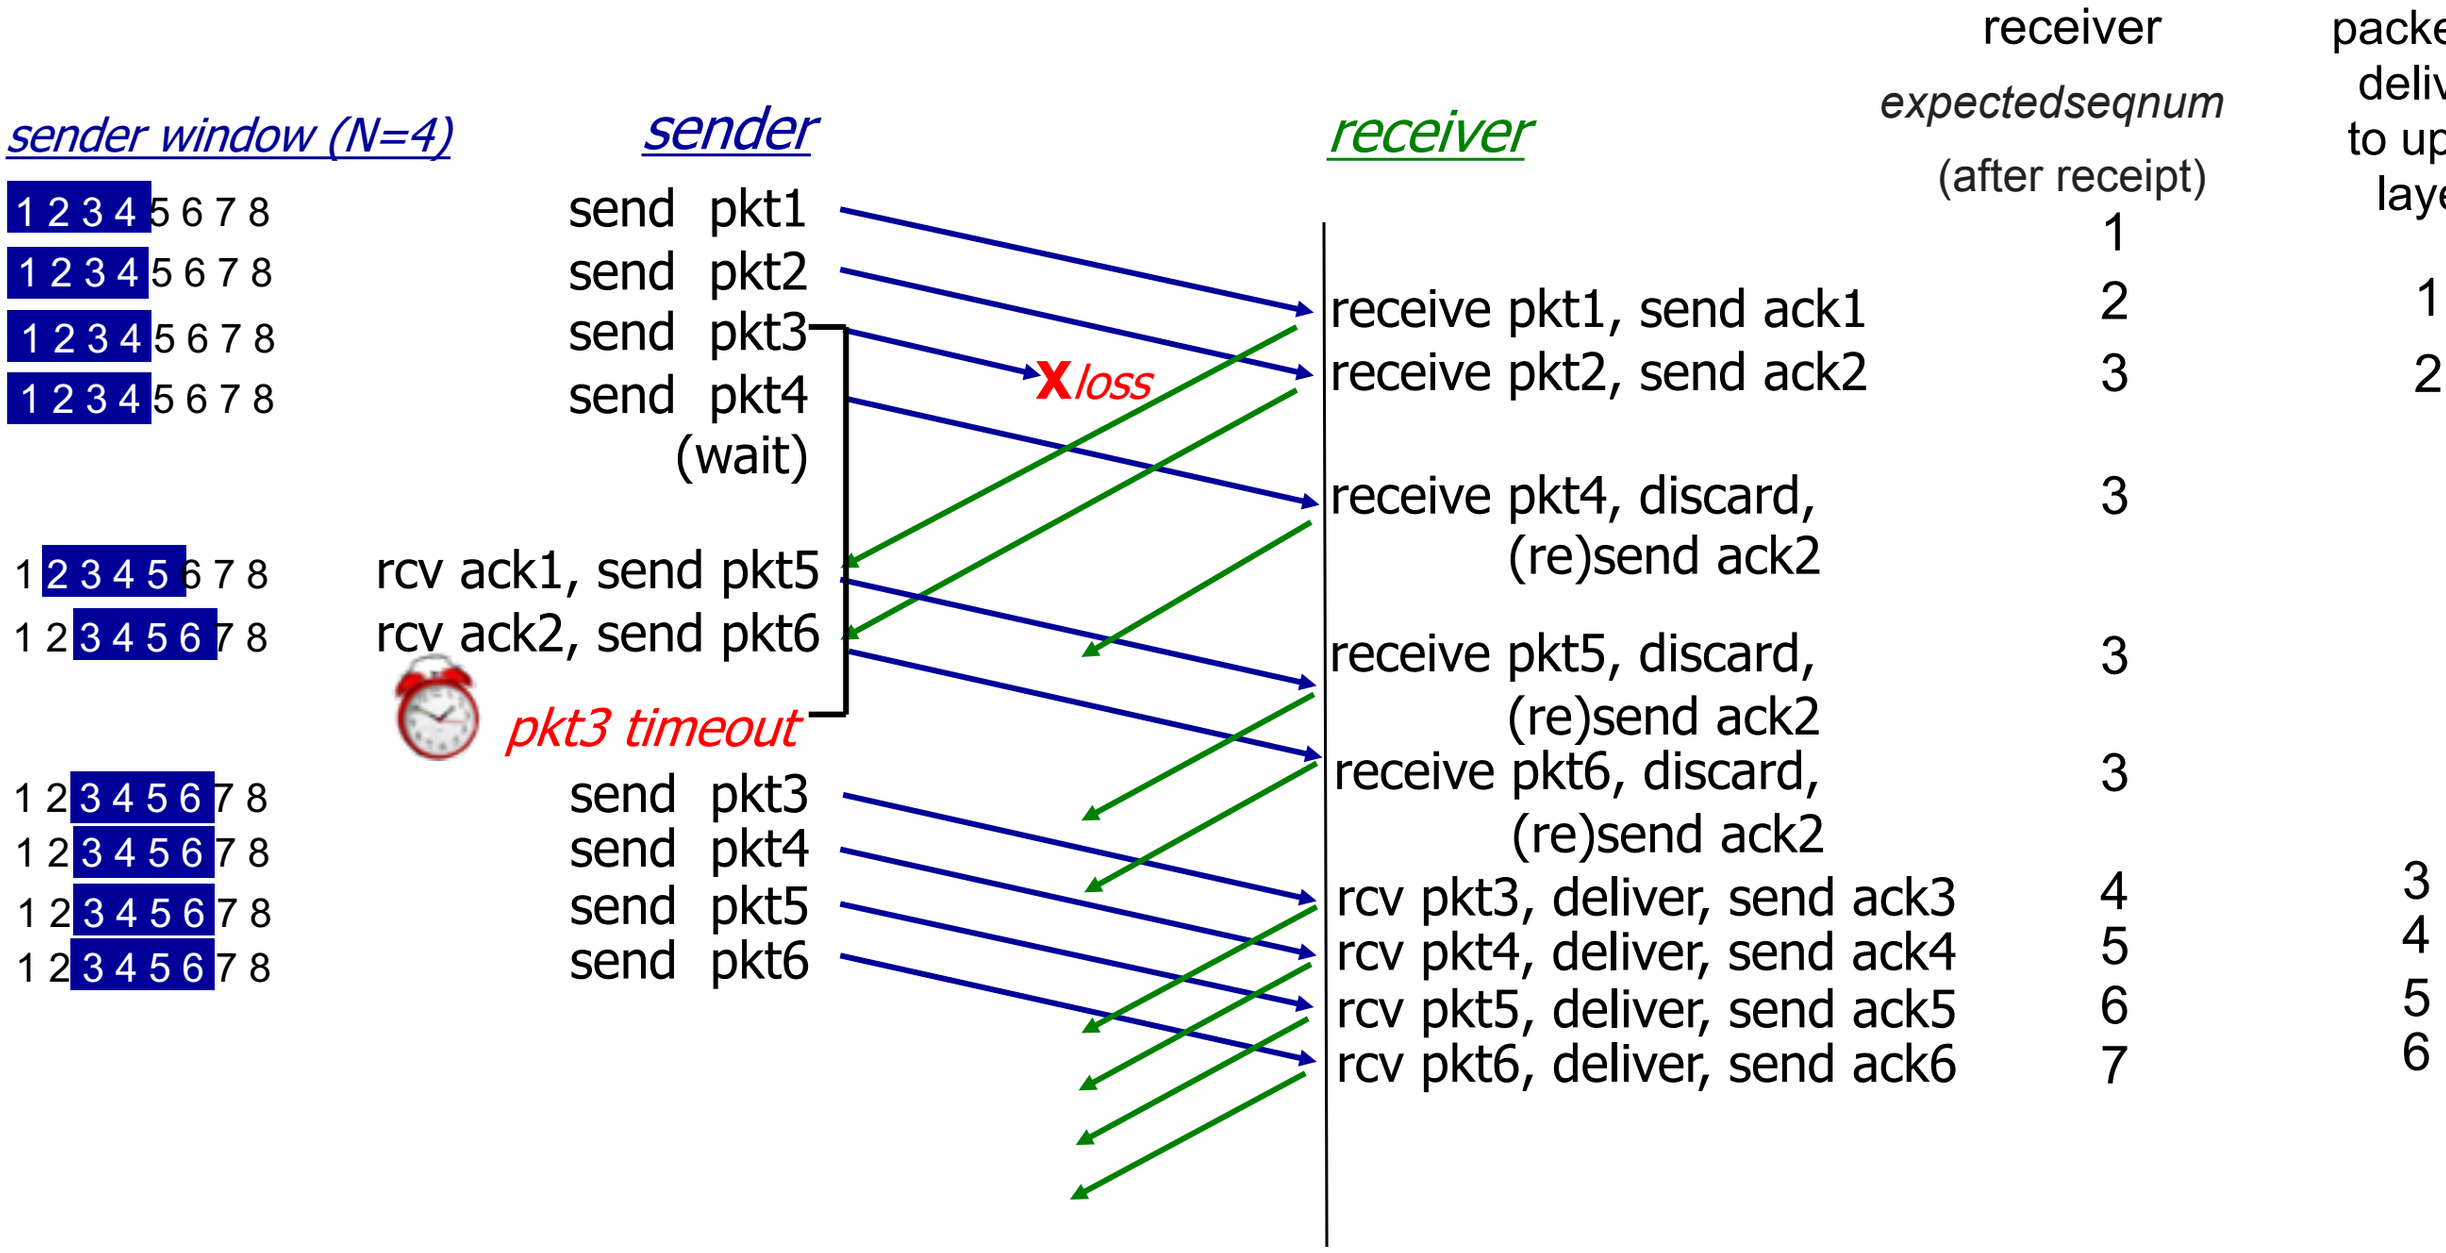

GBN in action (2)

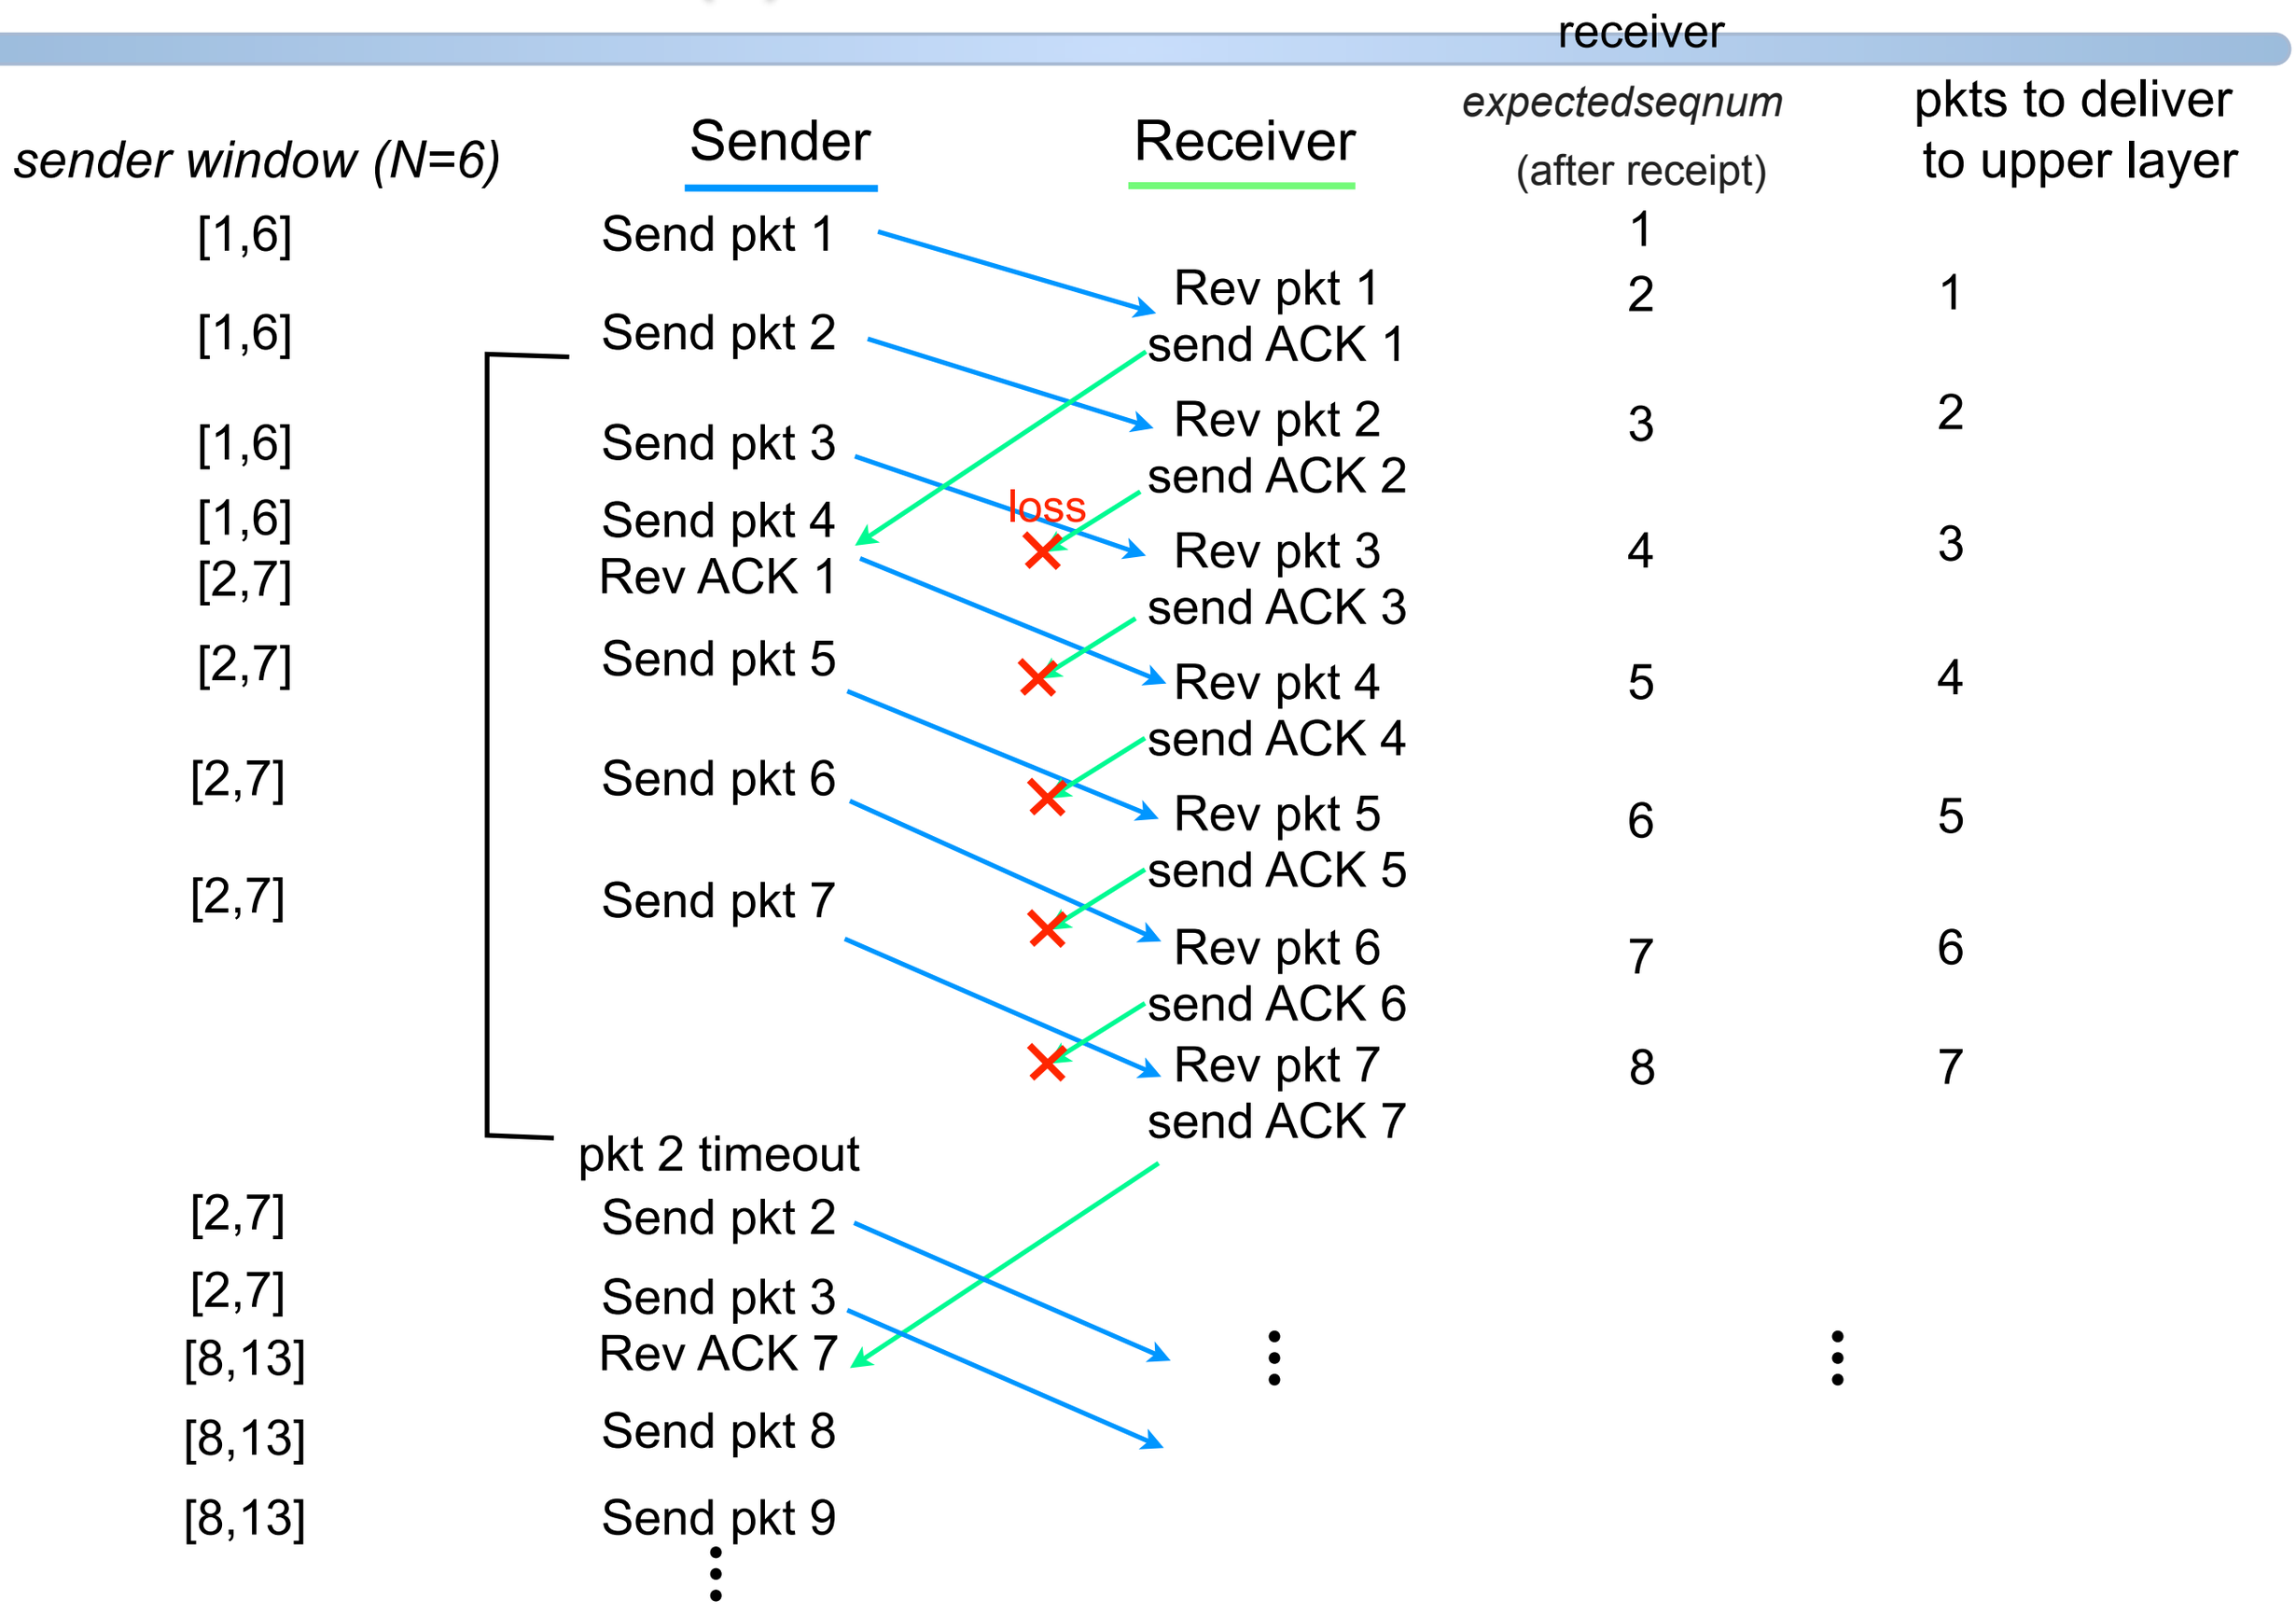

Go-Back-N (GBN): a major problem

- A major problem
  - many packets in pipeline when

    window size large and bandwidth-delay product large
  - single packet error causes retransmission of many packets

    not necessary

Selective Repeat

- Receiver individually acknowledges all correctly received packets
  - buffers packets as needed, for eventual in-order delivery to upper layer

- Sender only resends packets for which ACK not (correctly) received
  - sets timer for each unACKed packet
  - sender window
    - *N* consecutive seq #'s
    - limits seq #s of sent, unACKed pkts

Selective Repeat: sender/receiver windows

Sender window

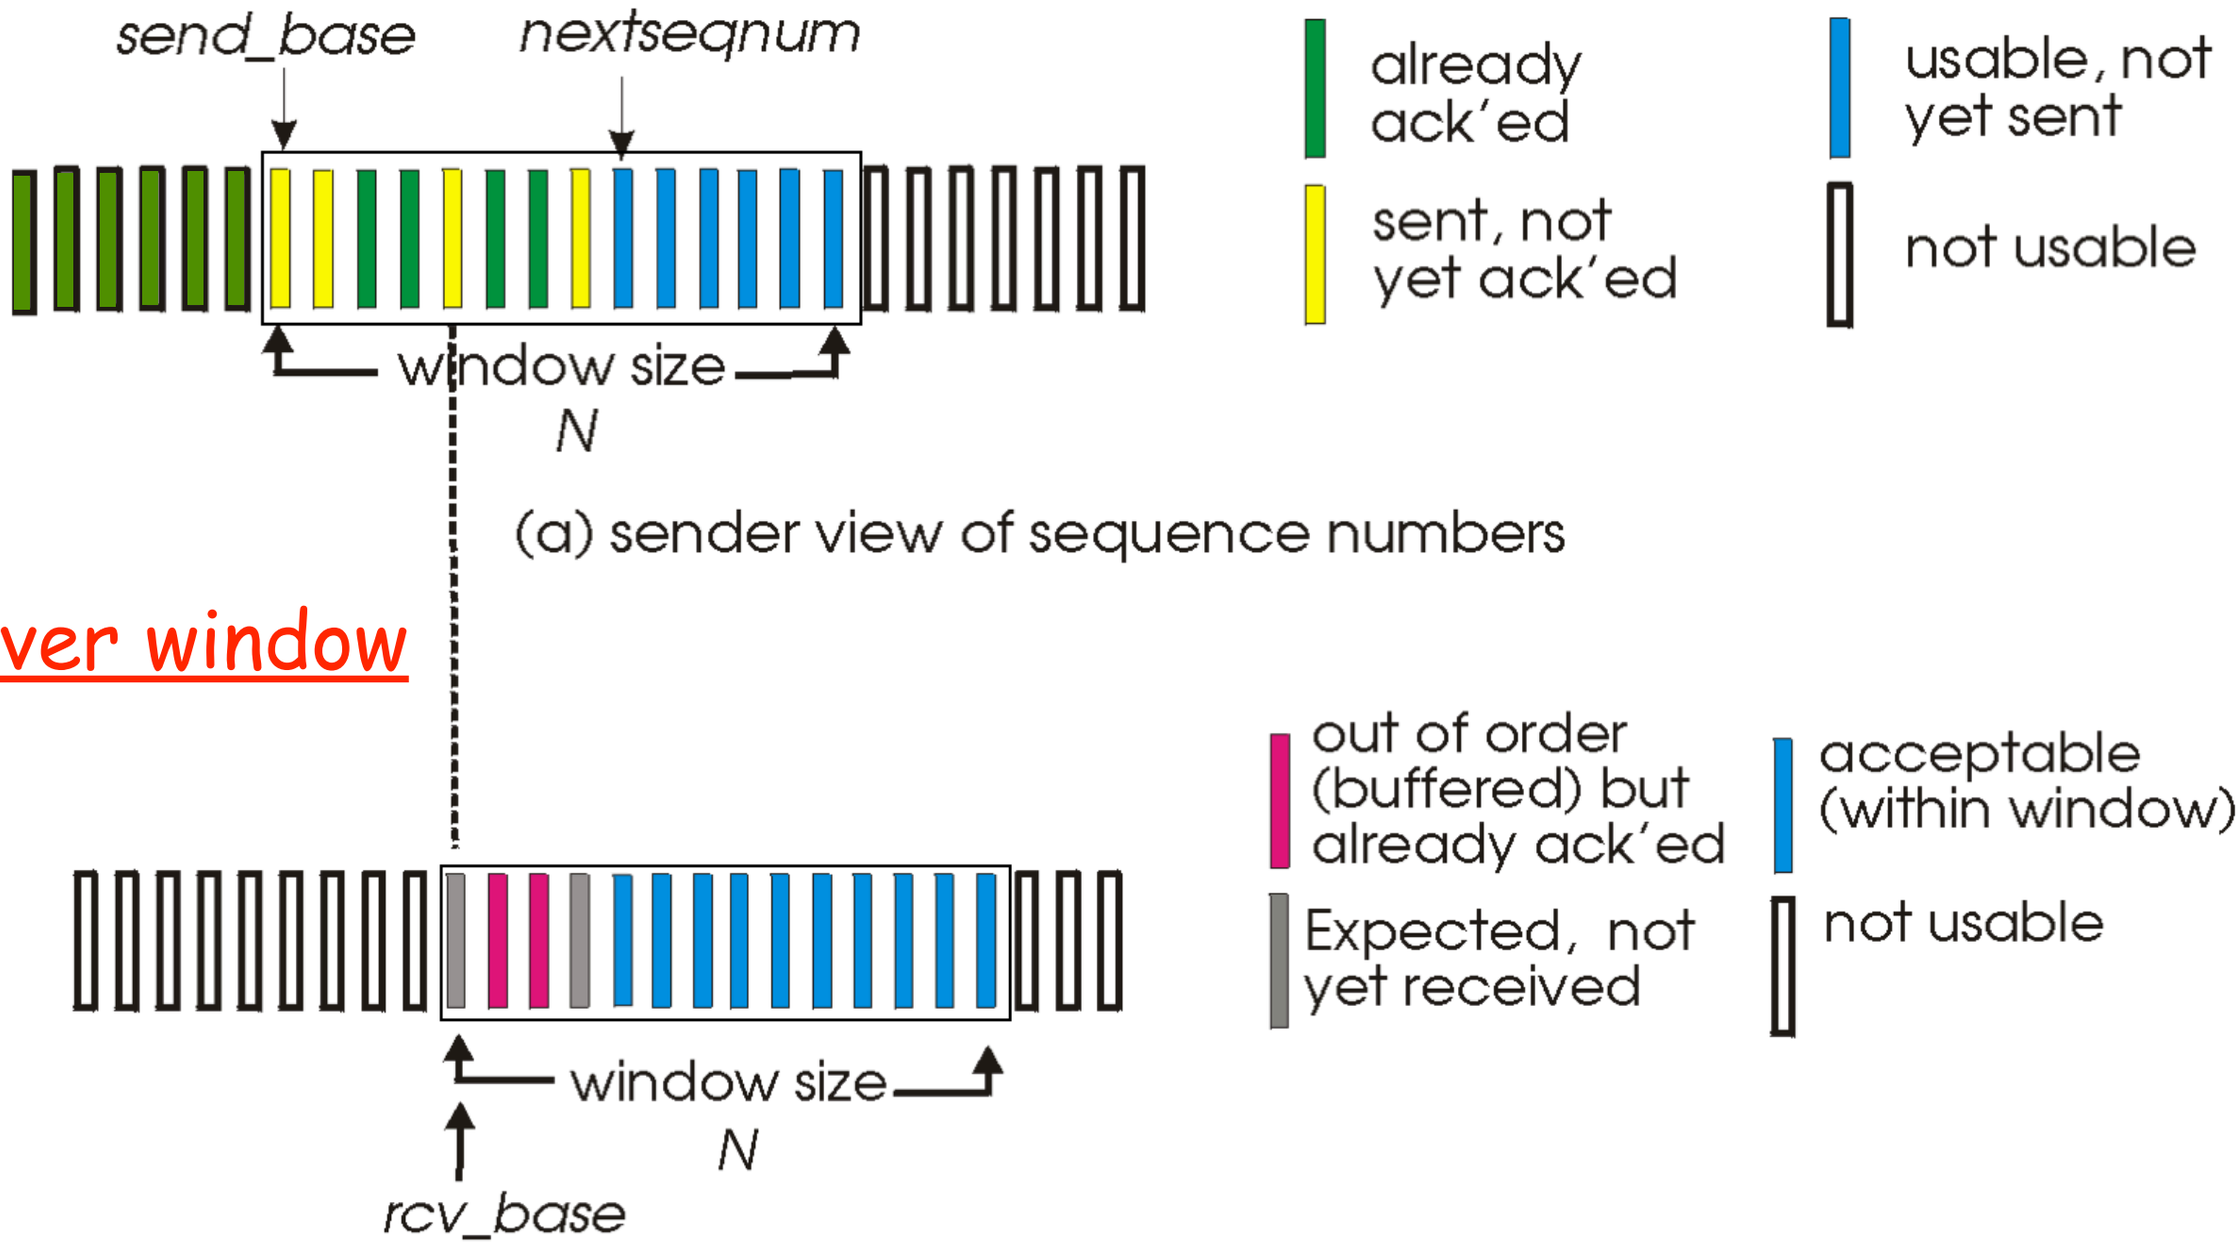

Receiver window

Selective Repeat: sender protocol

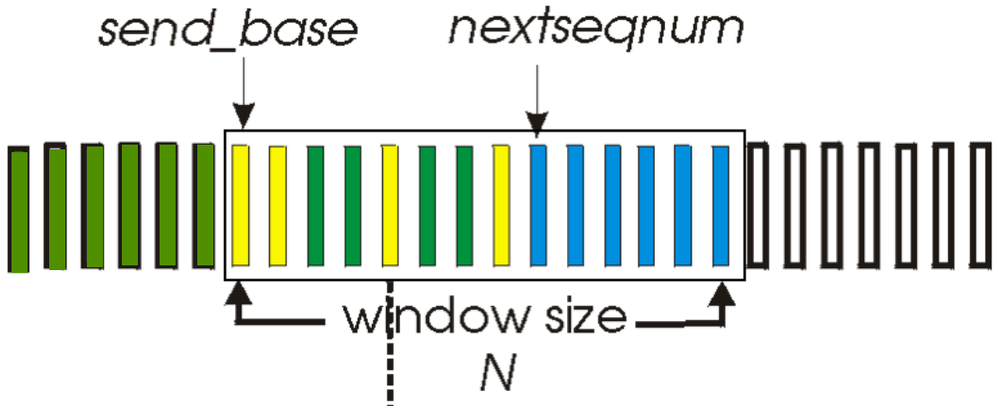

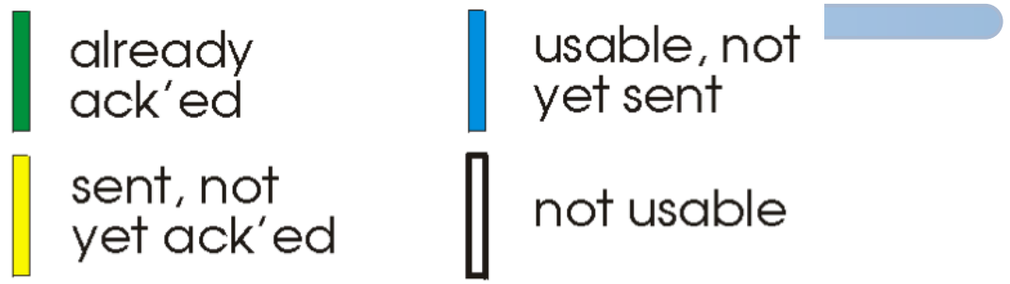

sender

data received from above:

- if next available seq # in window, send pkt and start timer on the pkt; otherwise, refuse data

timeout(n):

- resend pkt n, restart timer on n

ACK(n) in [send_base,send_base+N-1] correctly received:

- mark pkt n as received and stop its timer
- if n is smallest unACKed pkt, advance send_base to next unACKed seq #

Selective Repeat: receiver protocol

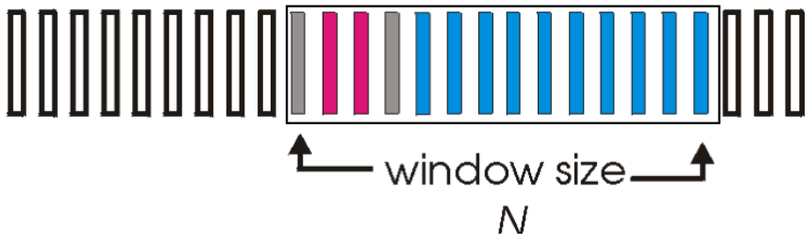

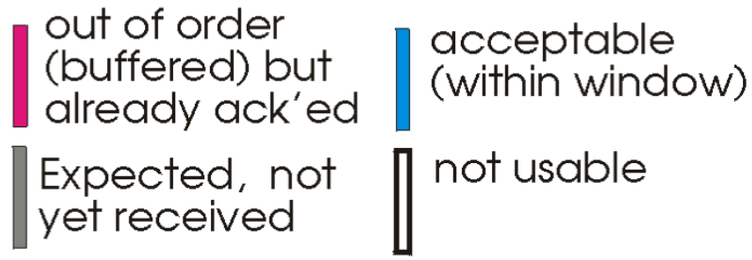

receiver

pkt n in [rcv_base, rcv_base+N-1]
correctly received:

- send ACK(n)
- out-of-order: buffer
- in-order: deliver (also deliver buffered, in-order pkts), advance rcv_base to next not-yet-received pkt

pkt n in [rcv_base-N,rcv_base-1]

- Send ACK(n)

otherwise:

- ignore

Why?

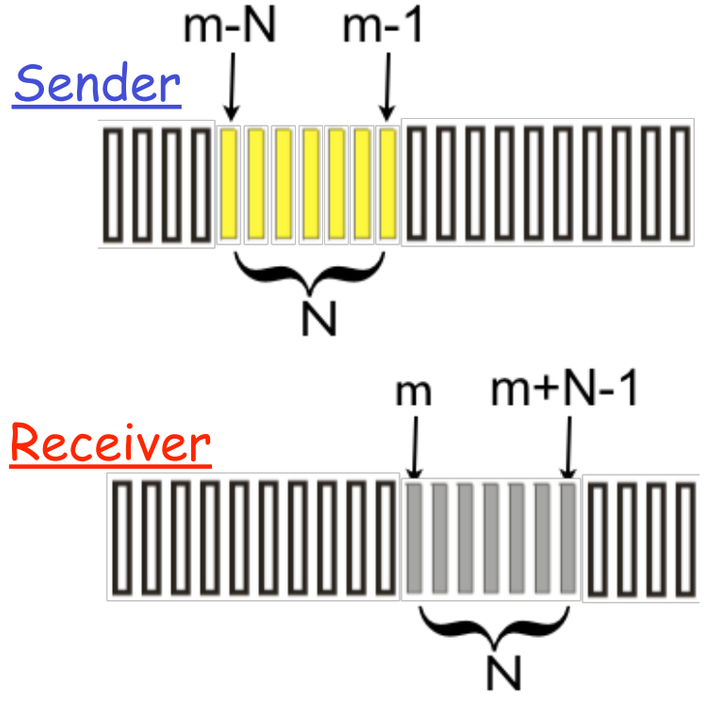

Selective Repeat in action (1)

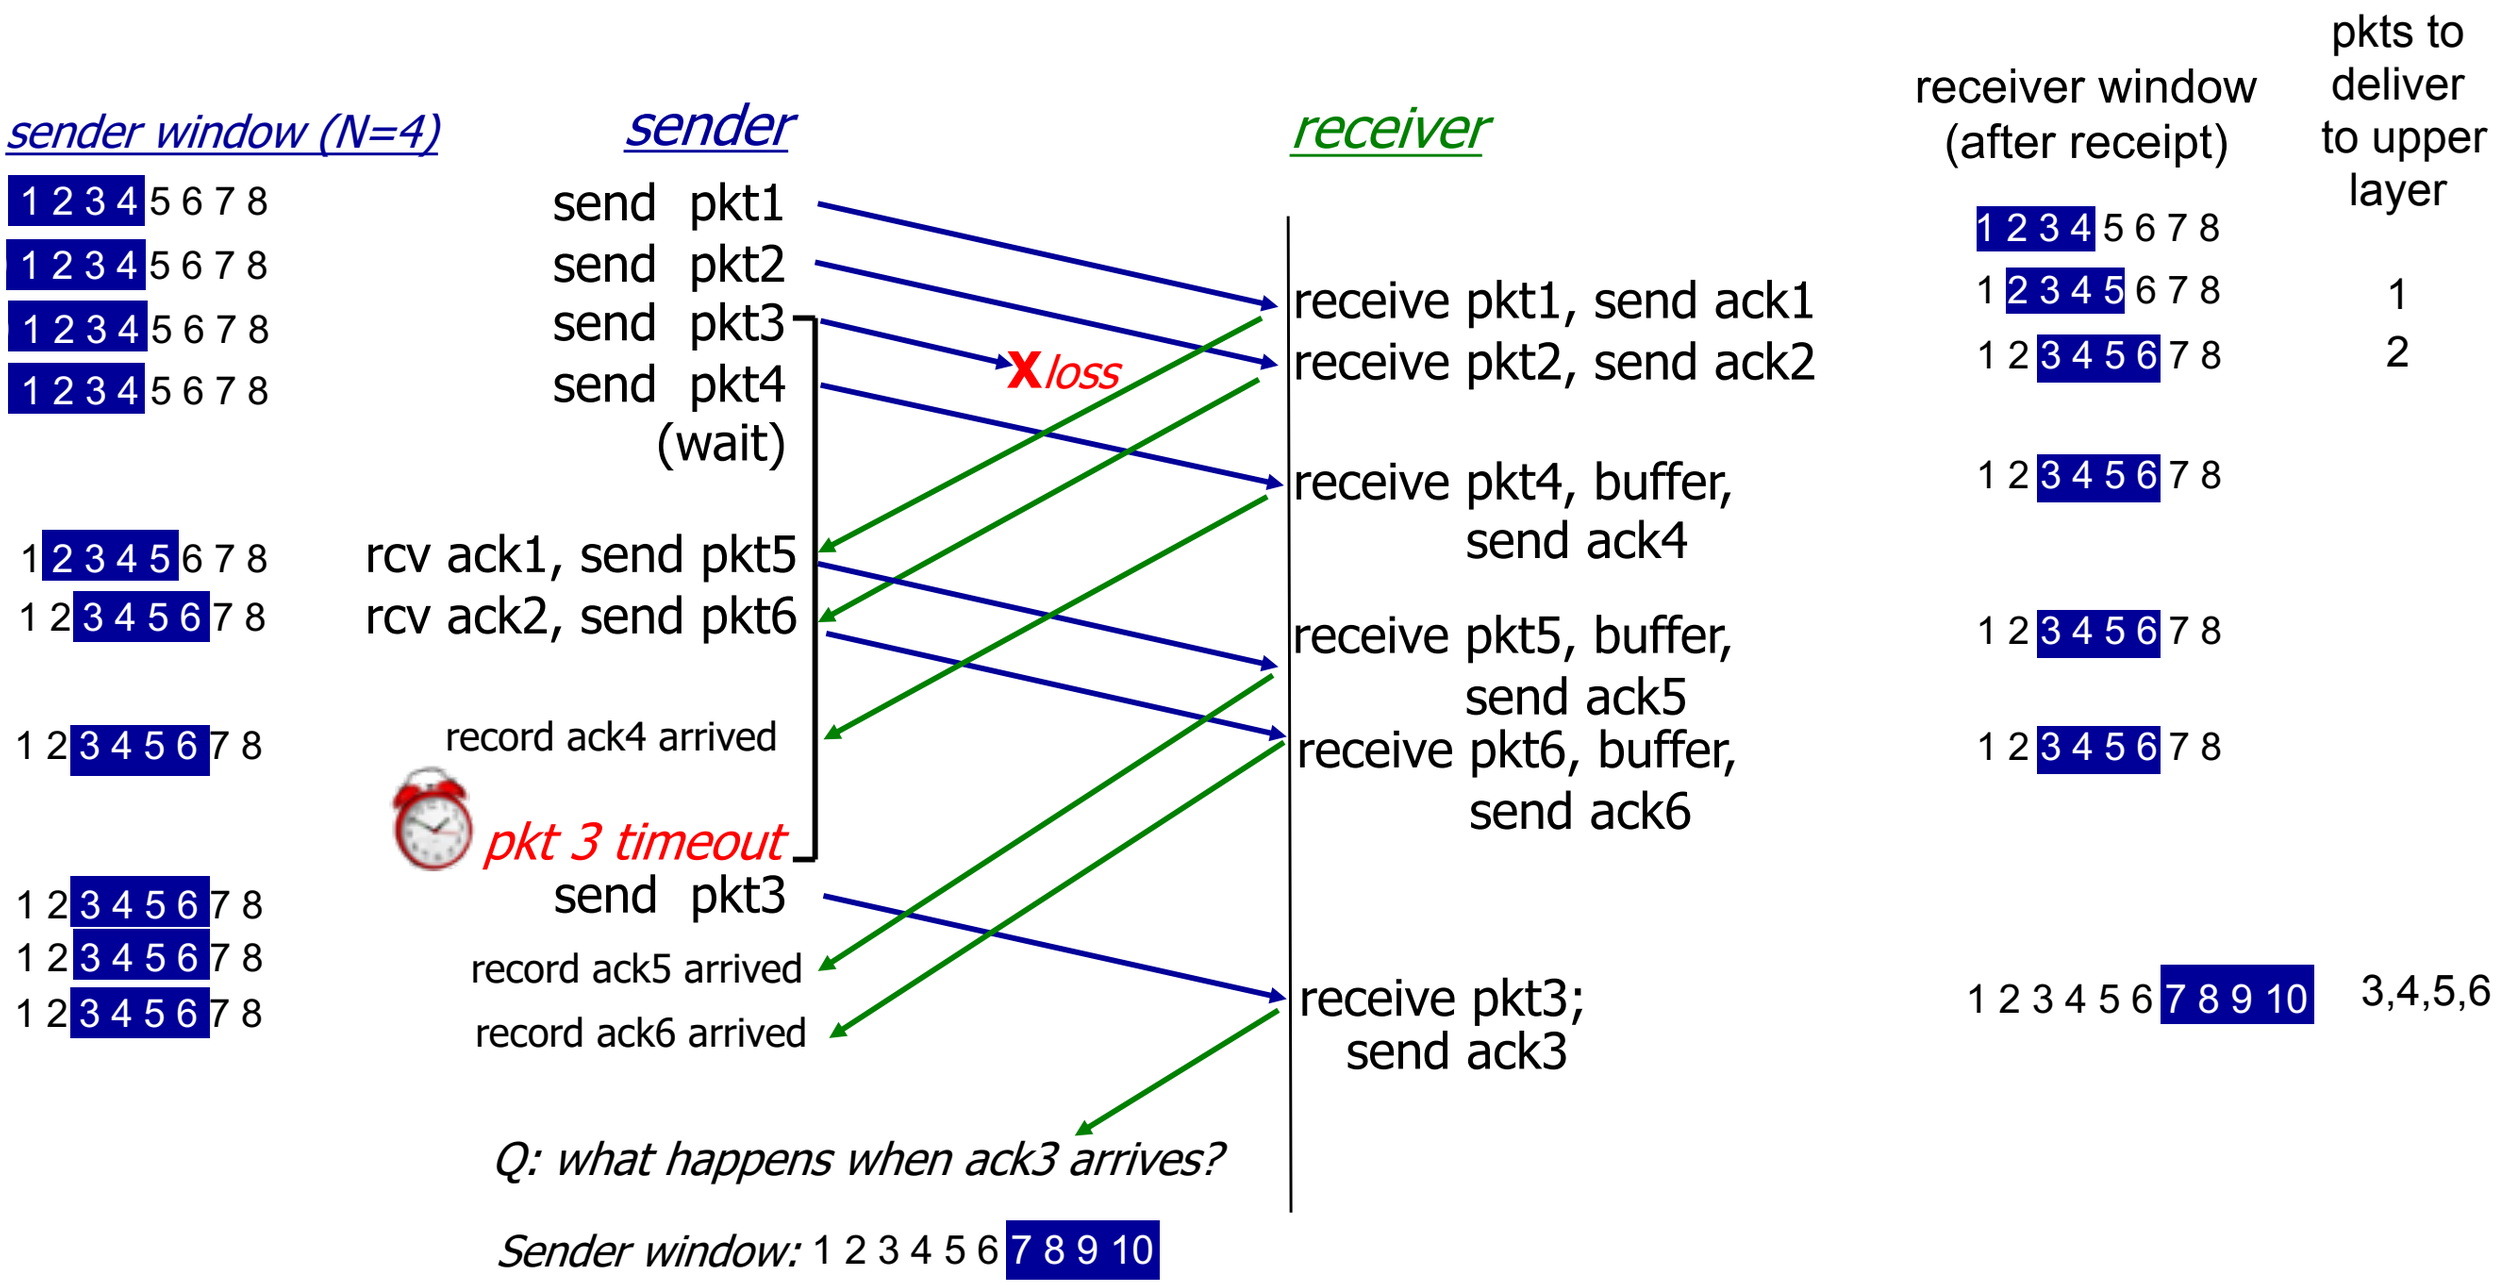

Selective Repeat in action (2)

N=6

ACKs in buffer	sender window	Sender	Receiver	receiver window (after receipt)	pkts in buffer	pkts to deliver to upper layer
	[1,6]	Send pkt 1		[1,6]		
	[1,6]	Send pkt 2 (loss)	Rev pkt 1 send ACK 1	[2,7]		1
	[1,6]	Send pkt 3				
	[1,6]	Send pkt 4	Rev pkt 3 send ACK 3	[2,7]	3	
	[2,7]	Rev ACK 1				
	[2,7]	Send pkt 5	Rev pkt 4 send ACK 4	[2,7]	3, 4	
3	[2,7]	Rev ACK 3	Rev pkt 5 send ACK 5	[2,7]	3,4,5	
	[2,7]	Send pkt 6	Rev pkt 6 send ACK 6	[2,7]	3,4,5,6	
	[2,7]	Send pkt 7	Rev pkt 7 send ACK 7	[2,7]	3,4,5,6,7	
		pkt 2 timeout				
	[2,7]	Send pkt 2	Rev pkt 2 send ACK 2	[8,13]		2,3,4,5,6,7
3, 7	[2,7]	Rev ACK 7				
7	[4,9]	Rev ACK 2 Send pkt 8				

Selective Repeat: window size vs. seq # size

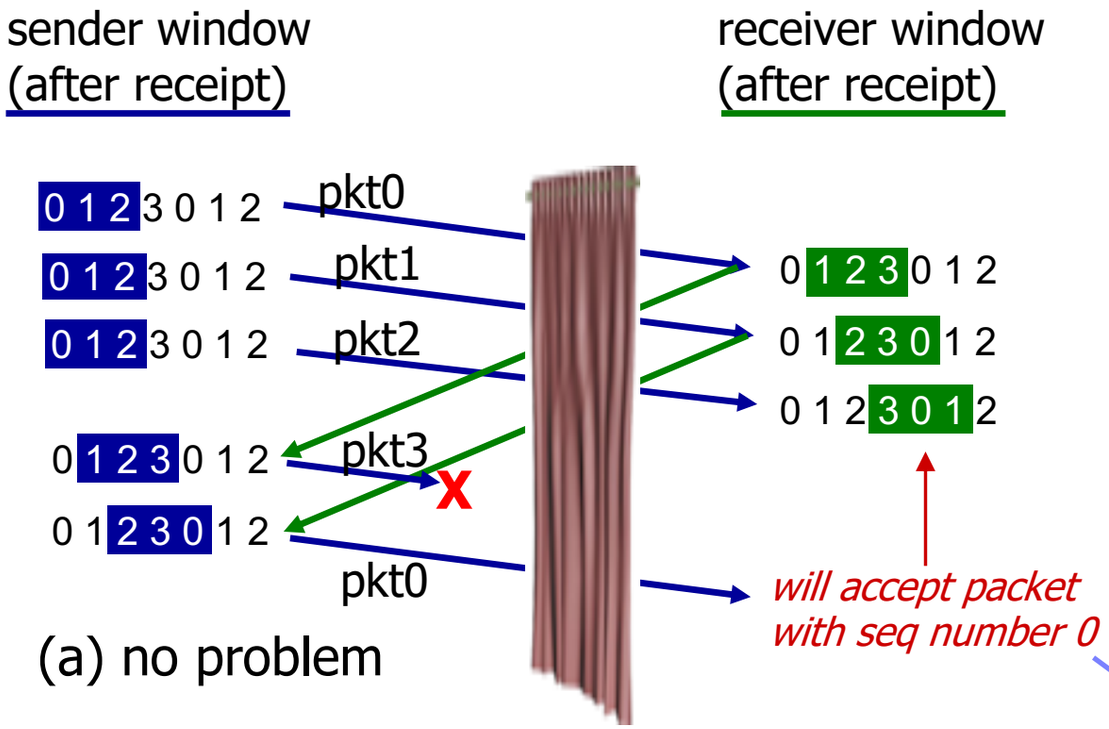

receiver can't see sender side.
receiver behavior identical in both cases!
something's (very) wrong!

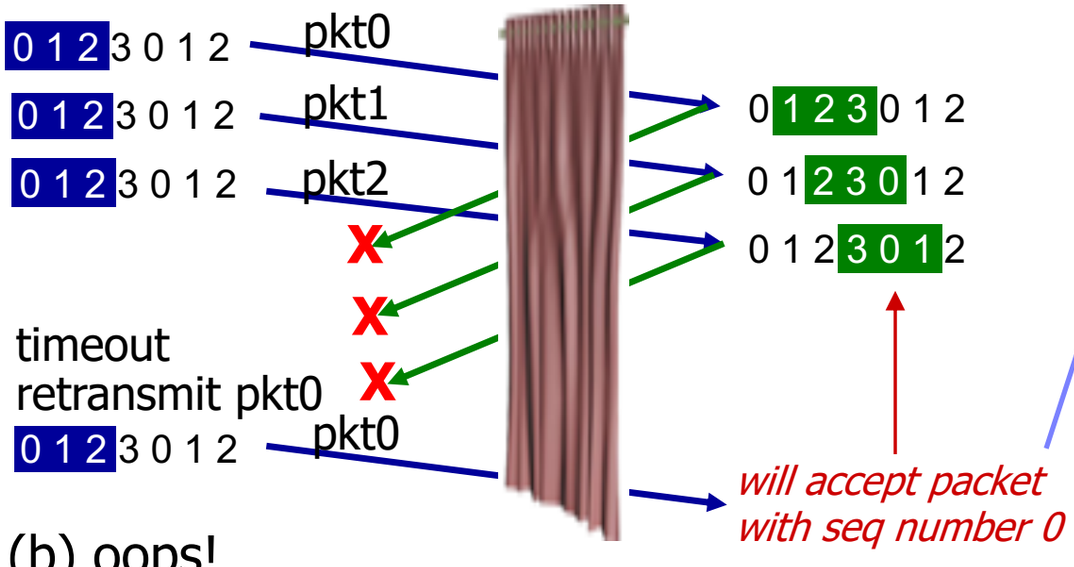

Q: what relationship between seq. # size and window size?

- Example
  - window size: N=3
  - seq #'s: 0,1,2,3

Receiver sees no difference in two scenarios!

in (b), incorrectly passes duplicate data as new

Selective Repeat: window size vs. seq # size (cont'd)

- Relationship between window size *N* and seq. # size *q* (2^k)
  - seq. # space must be large enough to fit the entire receiver window and the entire sender window
  - The extreme scenario

    receiver expects pkts [*m, m+N-1*]

    ACKs for pkt [*m-N,m-1*] are still propagating back

    sender window [*m-N,m-1*]

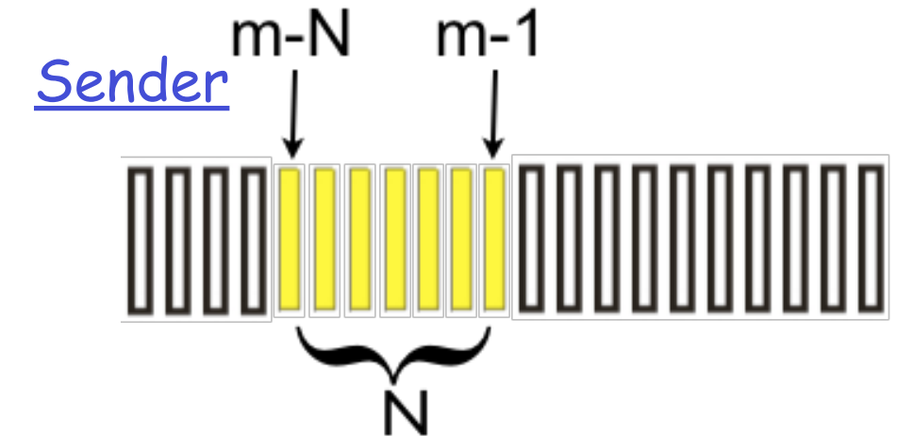

Therefore

Selective Repeat: $q \geq 2N$

Others

GBN: $q \geq N+1$

Stop-and-Wait: $q \geq 2$

- Required reading
  - Computer Networking: A Top-Down Approach (8th Edition)
    Ch 3.4.2, 3.4.3, 3.4.4
- Interactive animation of GBN:
  - https://media.pearsoncmg.com/ph/esm/ecs_kurose_compnetwork_8/cw/content/interactiveanimations/go-back-n-protocol/index.html
- Interactive animation of SR:
  - https://media.pearsoncmg.com/ph/esm/ecs_kurose_compnetwork_8/cw/content/interactiveanimations/selective-repeat-protocol/index.html
- Acknowledgement:
  - Some materials are extracted from the slides created by Prof. Jim F. Kurose and Prof. Keith W. Ross for the textbook.