



COMP 3234B

Computer and Communication Networks

2nd semester 2023-2024

Transport Layer (III)

Prof. C Wu

**Department of Computer Science
The University of Hong Kong**

Roadmap

Transport layer

■ Principles behind transport-layer services

multiplexing/demultiplexing (ILO1, 2)

reliable data transfer:(ILO 2, 3)

rdt 1.0

rdt 2.0, 2.1, 2.2

rdt 3.0

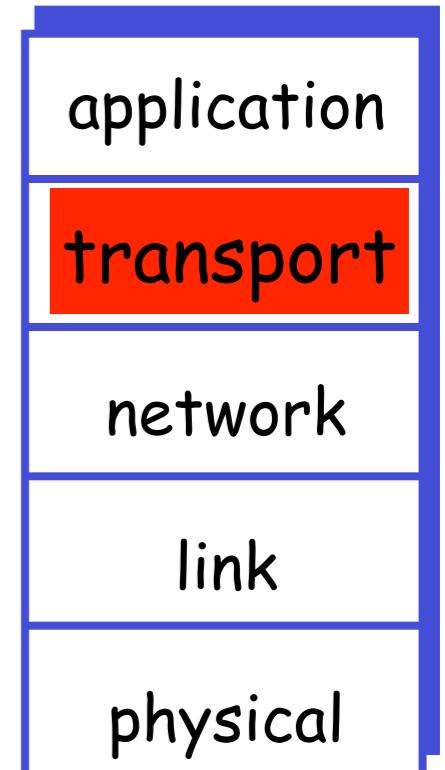
GBN

selective repeat

■ Transport protocols in the Internet (ILO 2, 3)

TCP: RDT, flow control, congestion control

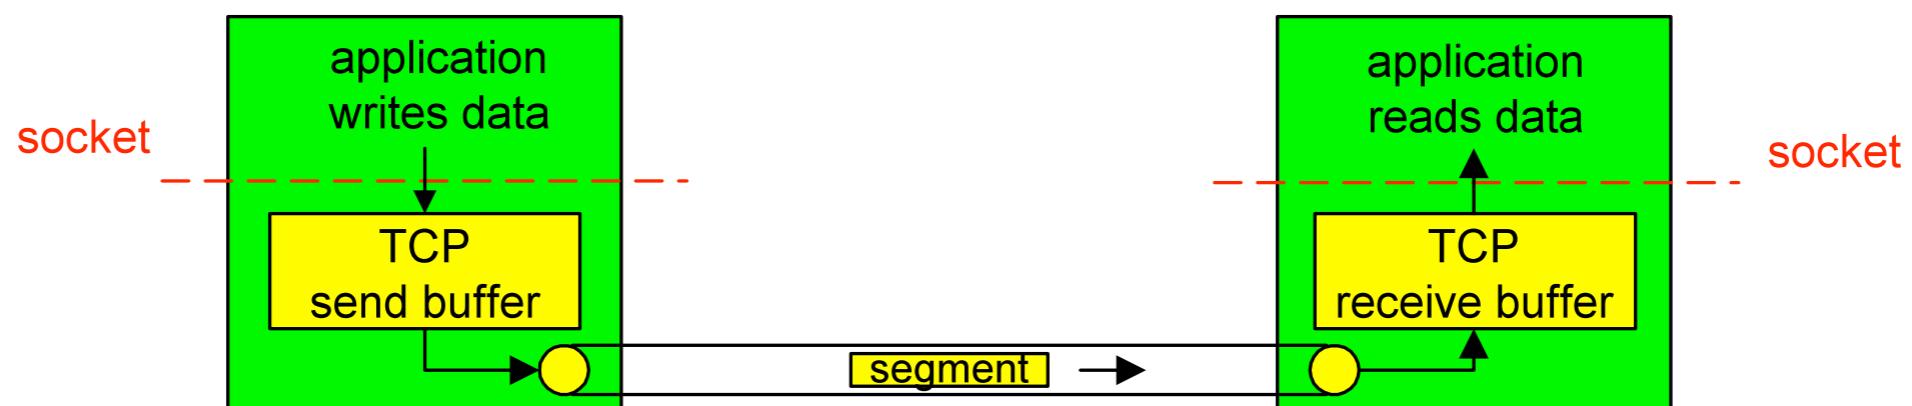
UDP



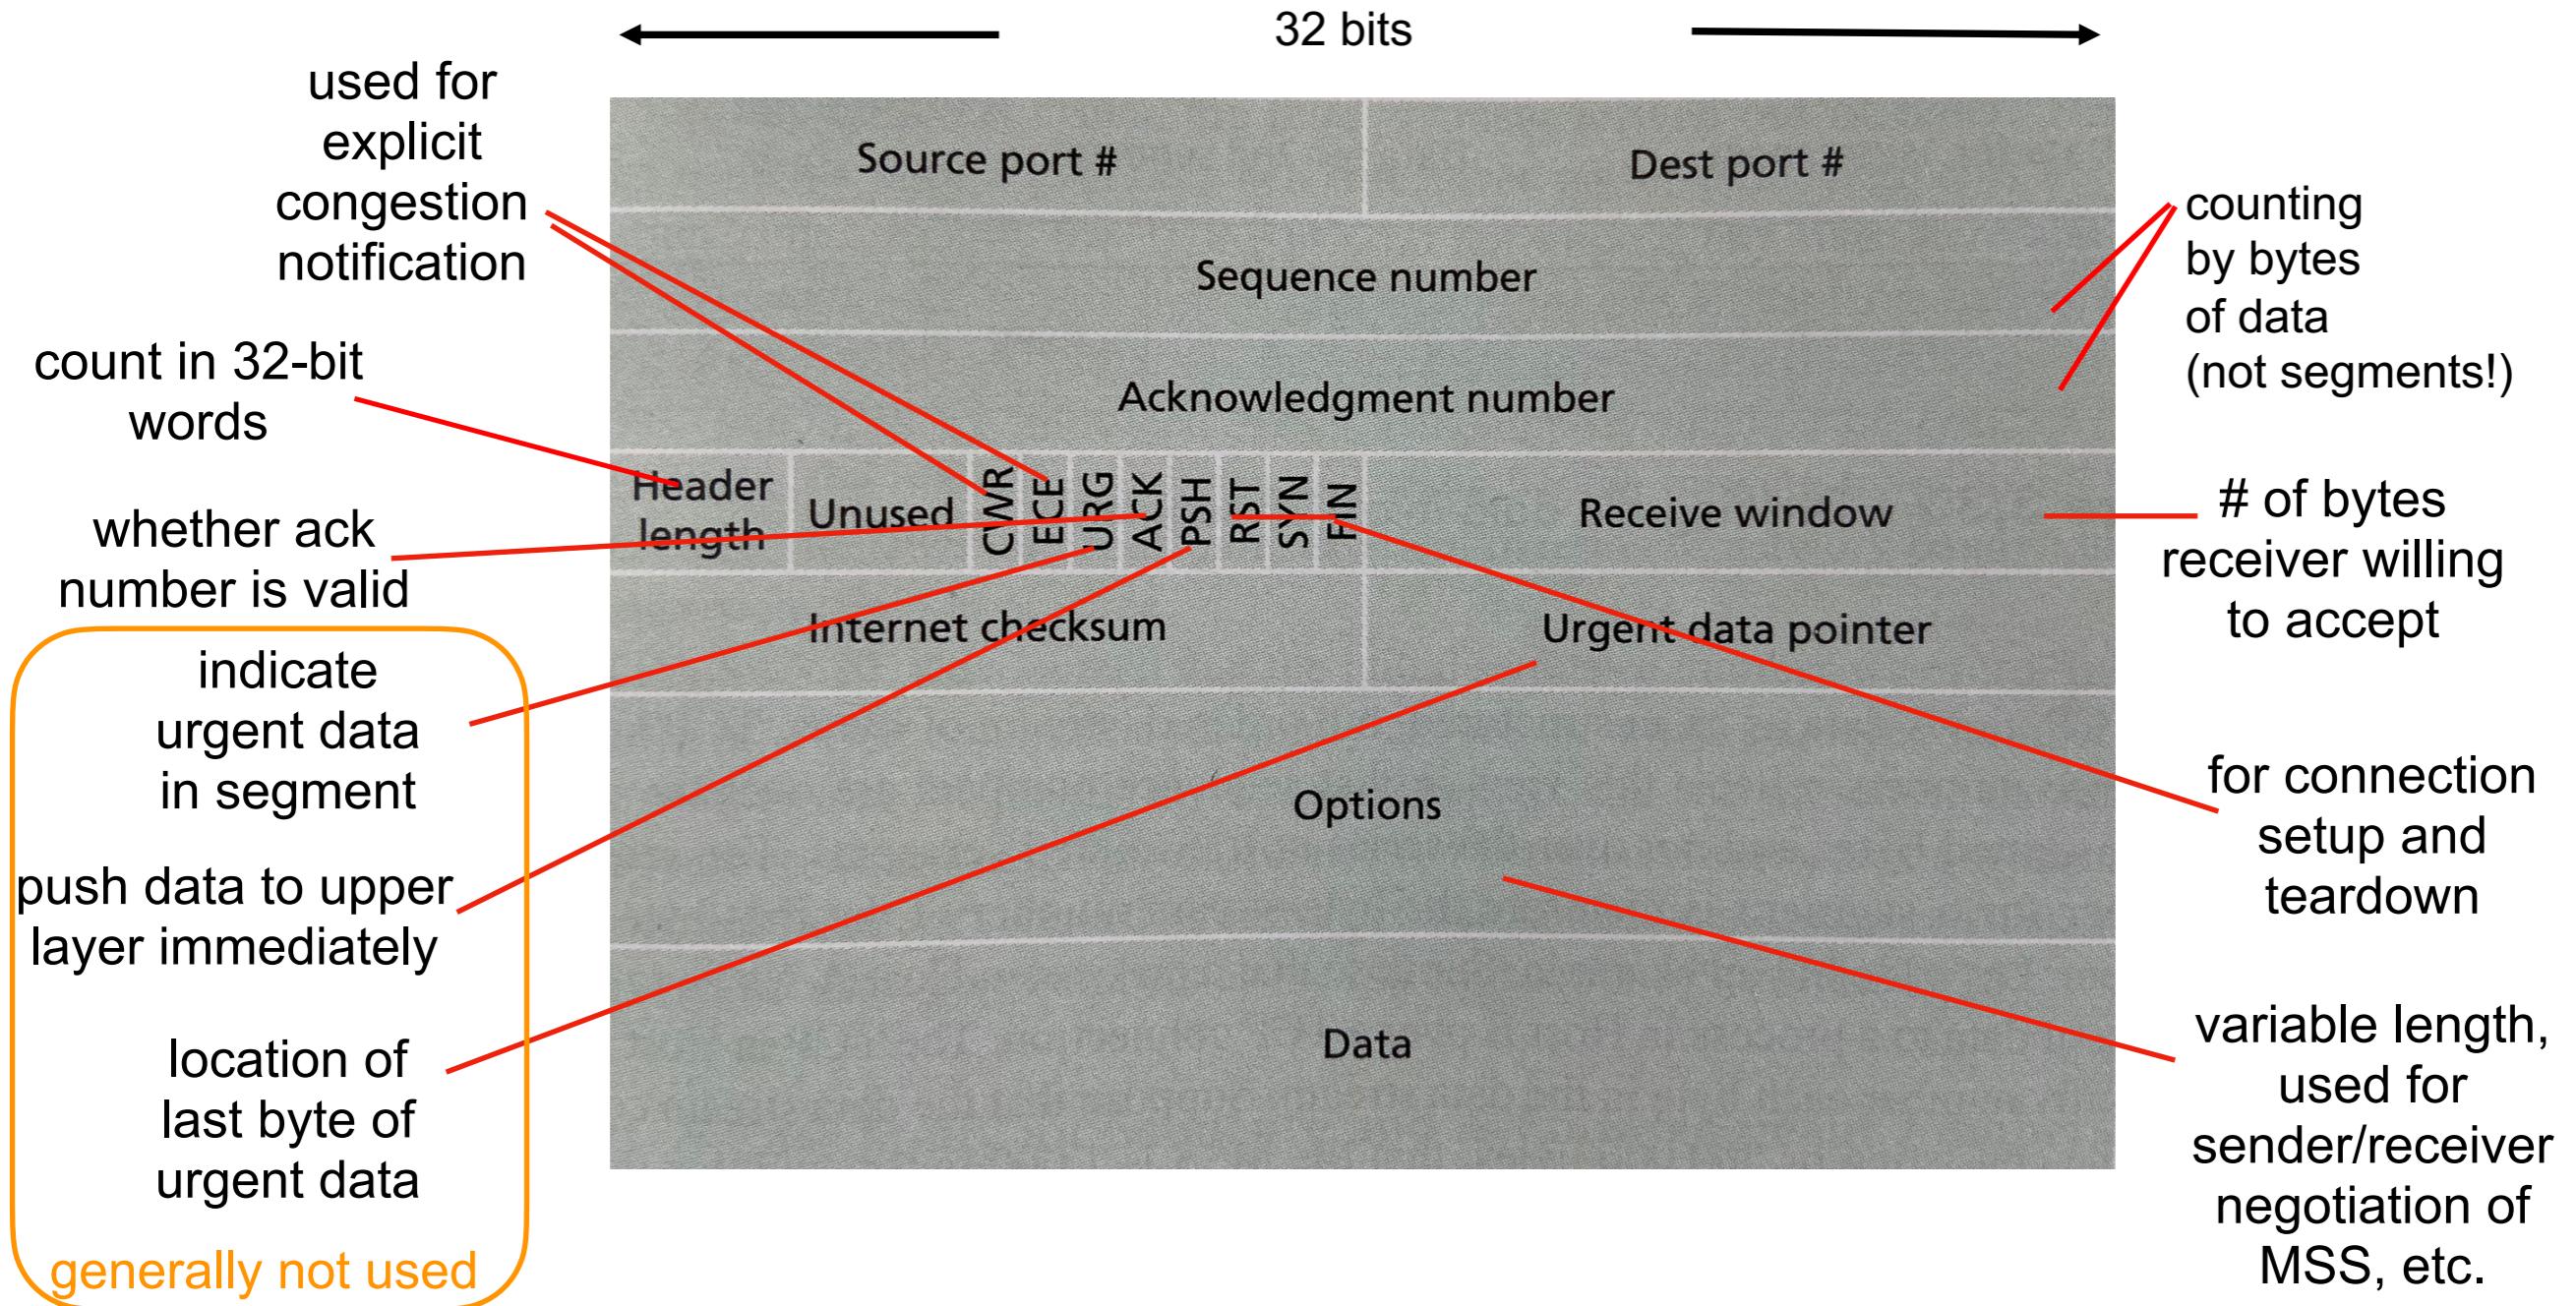
TCP: overview

RFCs: 793, 1122, 1323, 2018, 2581

- Point-to-point
 - one sender, one receiver
- Reliable data transfer
 - in-order byte stream
- Pipelined transfer
 - TCP congestion control and flow control (set window size)
- Full duplex service
 - bi-directional data flow in same connection
- Connection-oriented
 - handshaking (exchange of control msgs) before data exchange
- Flow and congestion controlled
 - sender will not overwhelm receiver
 - sender adjusts sending rate upon detecting congestion



TCP: segment structure

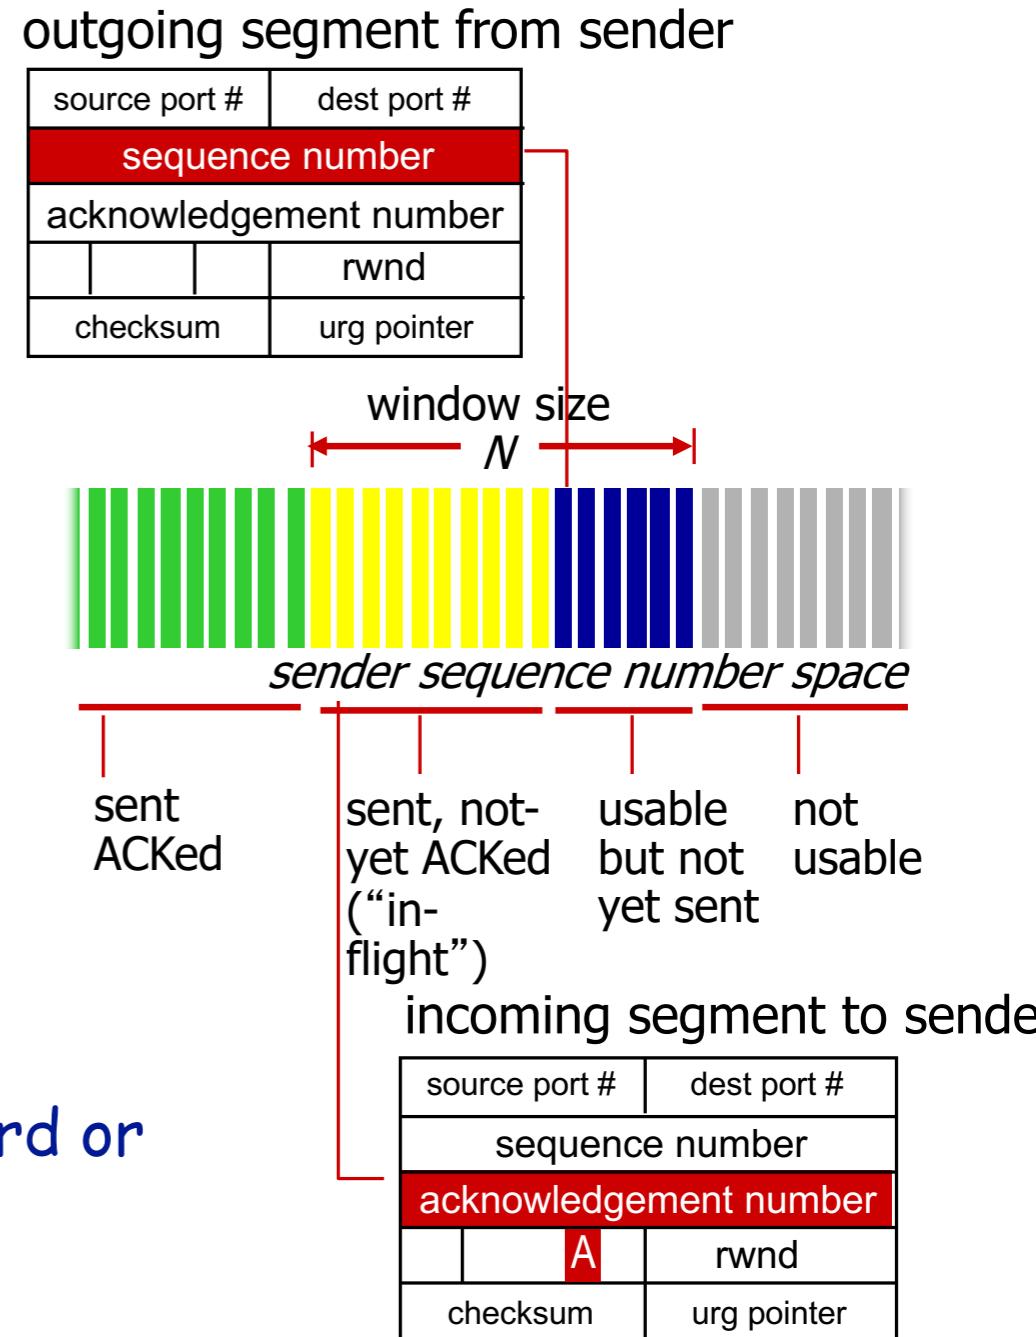


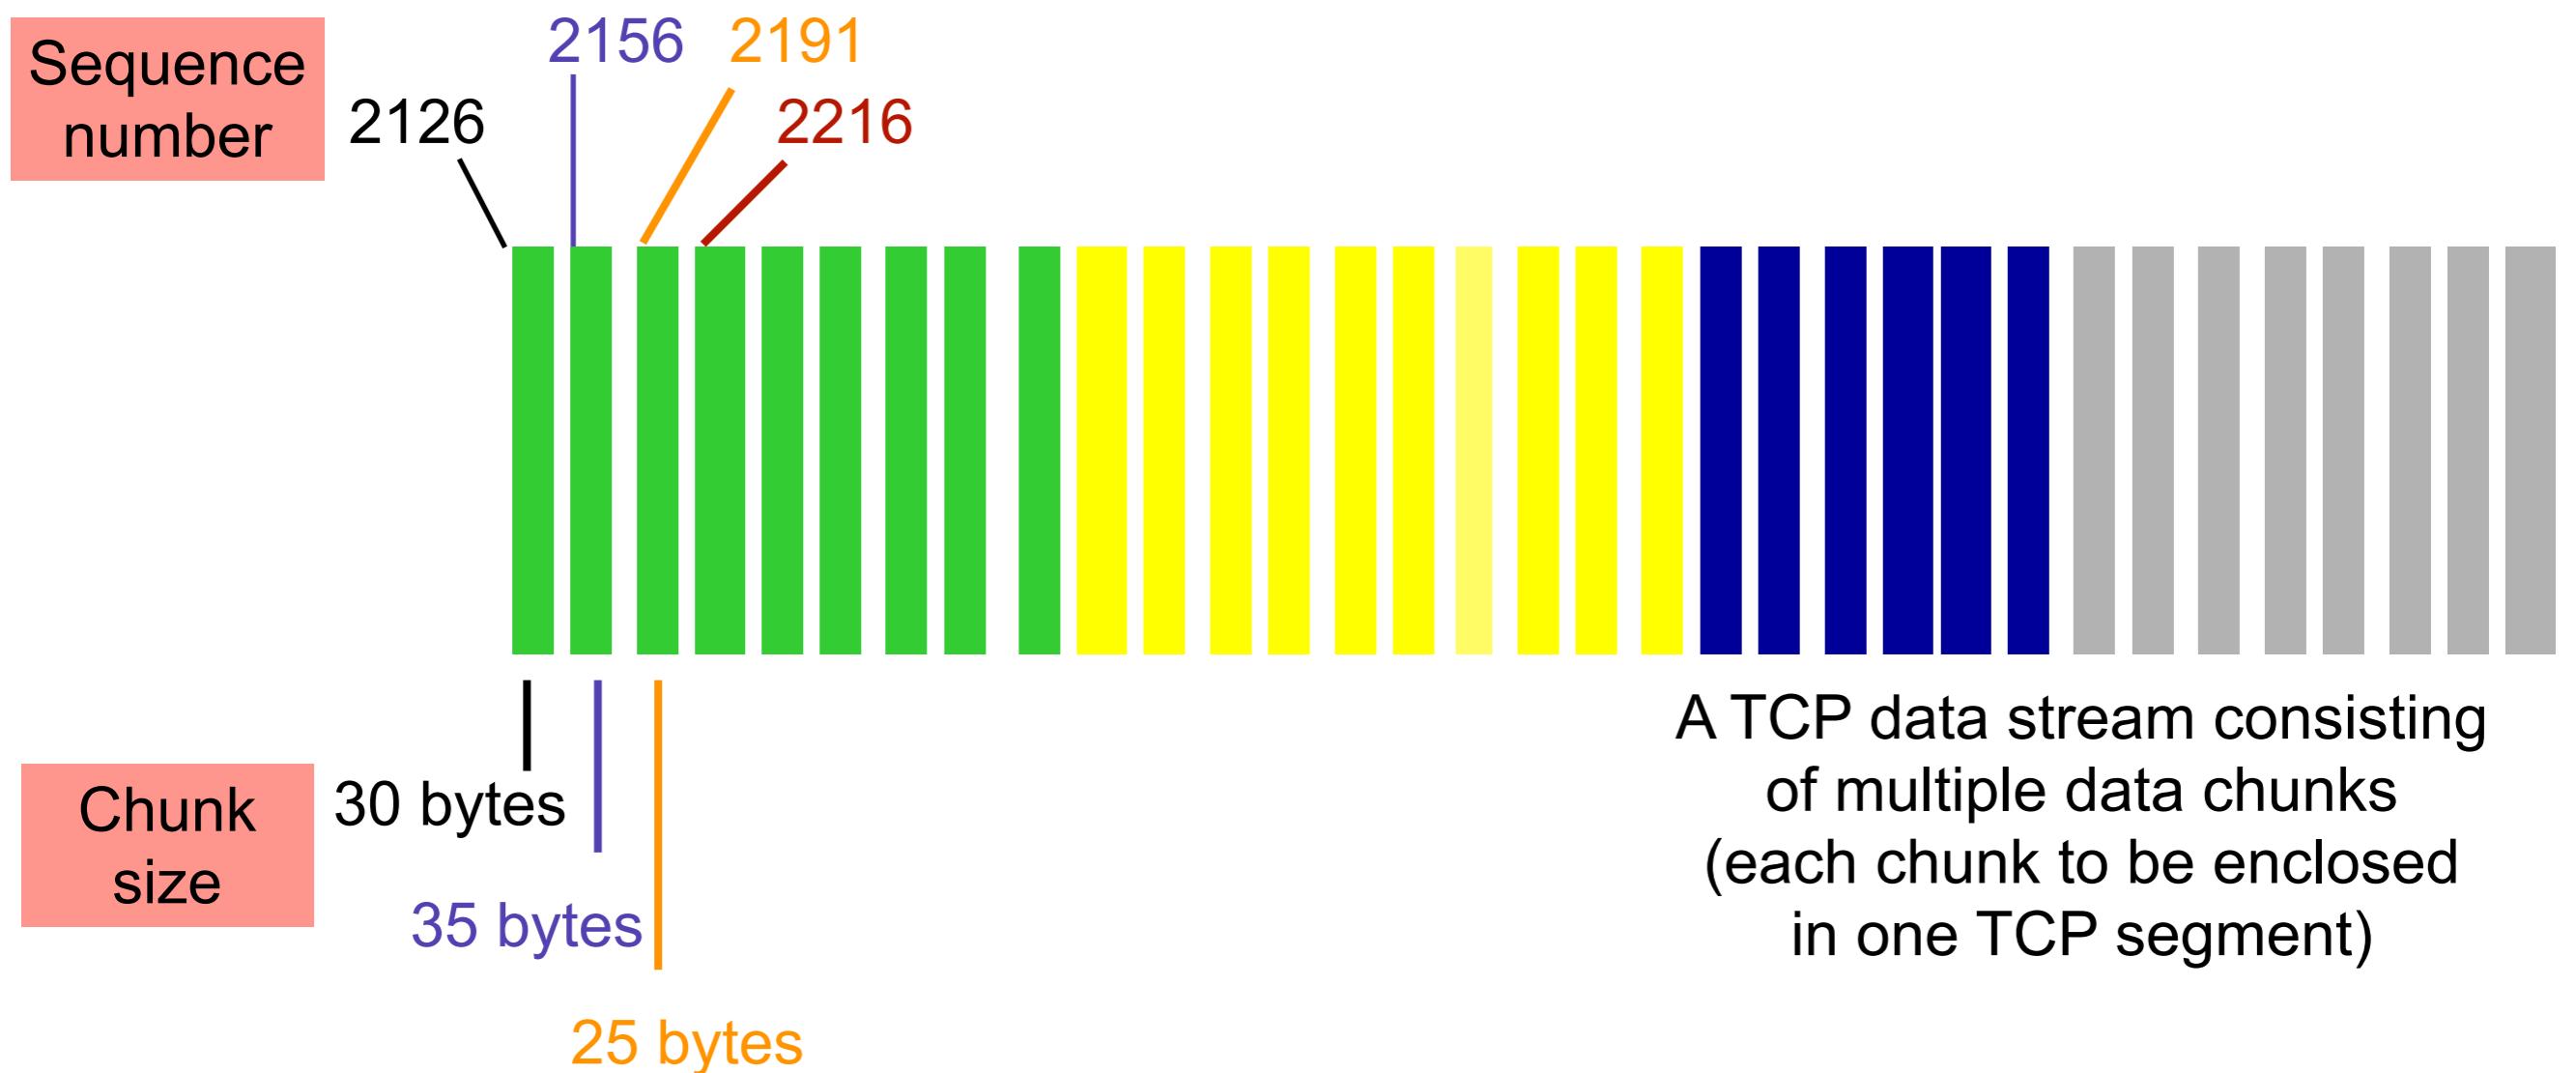
MSS (maximum segment size)

The maximum amount of application-layer data in a segment

TCP: seq. #'s and ACKs

- TCP views data as ordered stream of bytes
- Seq. #
 - byte-stream number of first byte in segment's data
not number of segments
 - initial seq. # is randomly chosen
- ACK #
 - seq. # of next byte expected from the other side
 - cumulative ACK
acknowledges bytes up to the first missing byte
 - ACK piggybacked in data segments
- How does receiver handle out-of-order bytes: discard or buffer?
 - TCP RFCs do not specify
up to implementor
in practice, buffer (as in Selective Repeat)

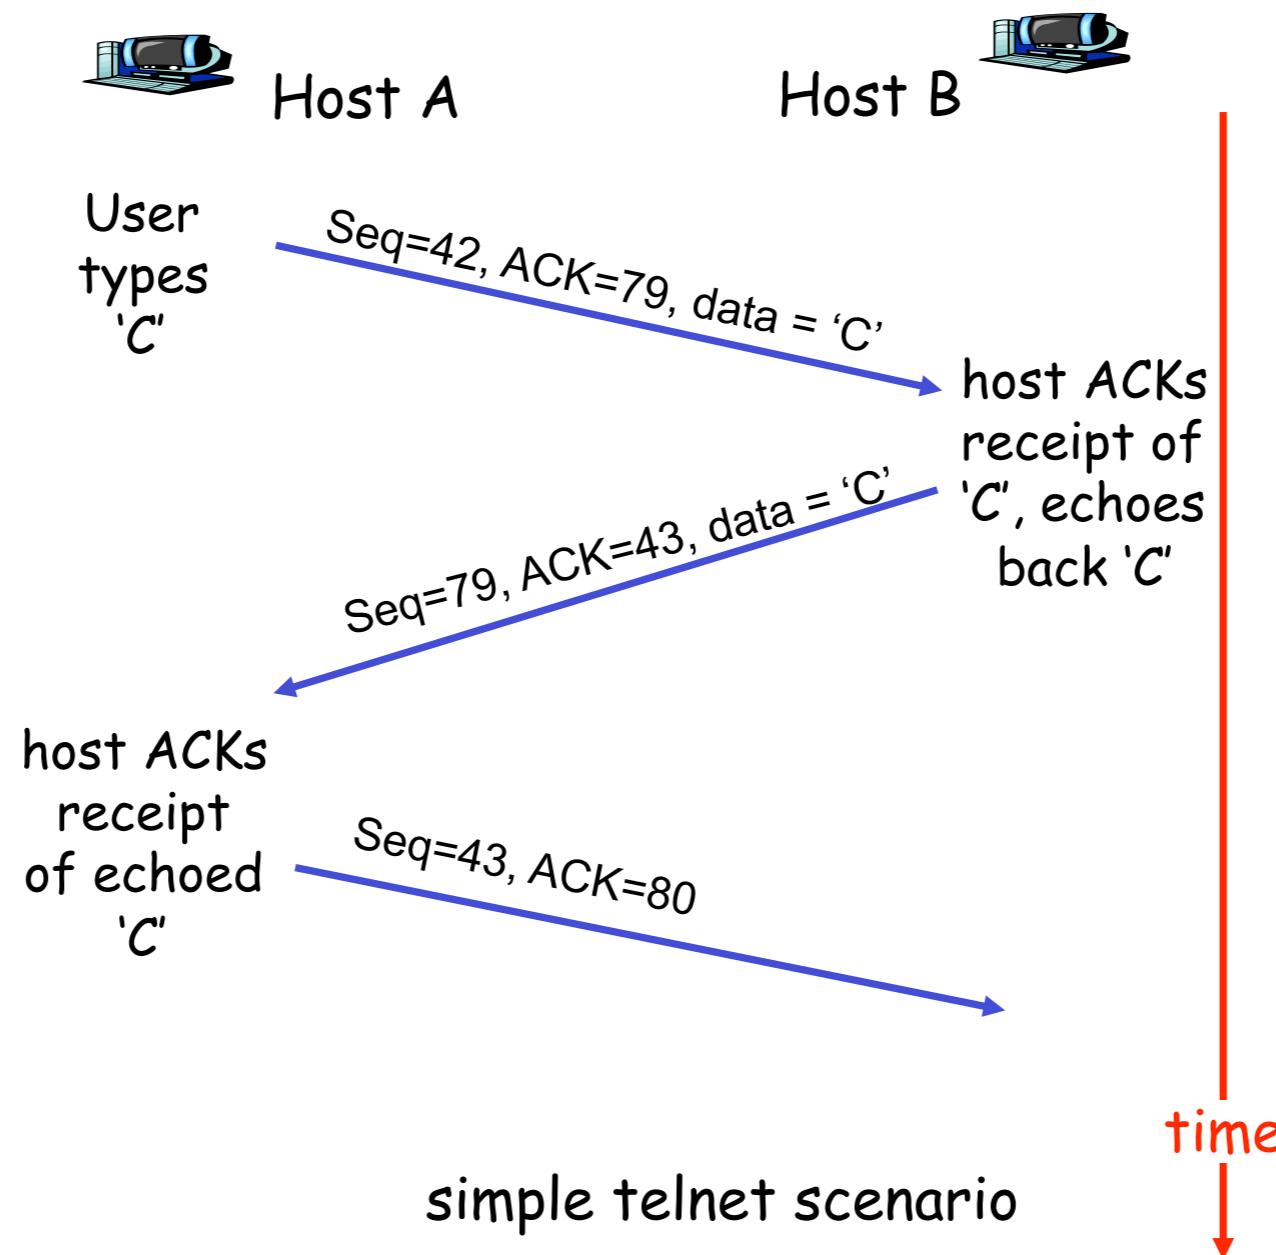




TCP: seq. #'s and ACKs (an example)

Telnet

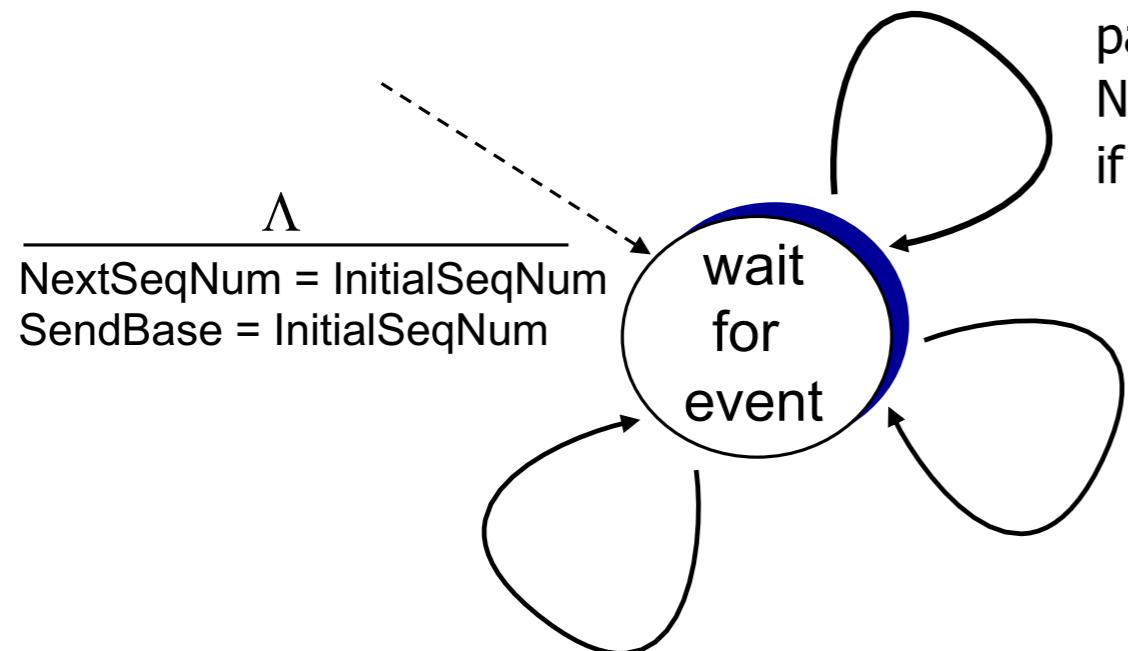
- application-layer protocol for remote login
- runs over TCP



TCP reliable data transfer

- TCP creates **reliable data transfer** service on top of IP's unreliable best-effort service
- Selective repeat or go-back-N? A hybrid.
 - Pipelined segments
 - Cumulative acks
 - Uses single retransmission timer
 - Retransmissions are triggered by:
 - timeout events
 - duplicate acks

TCP RDT: simplified sender



ACK received, with ACK field value y

```
if (y > SendBase) {  
    SendBase = y  
    /* SendBase-1: last cumulatively ACKed byte */  
    if (there are currently not-yet-acked segments)  
        start timer  
    else stop timer  
}
```

data received from app layer

- create segment with seq. #
- start timer if not already running
timer is for oldest unacked segment

data received from application above

create segment, seq. #: NextSeqNum

pass segment to IP (i.e., “send”)

NextSeqNum = NextSeqNum + length(data)

if (timer currently not running)

start timer

timeout

retransmit not-yet-acked segment
with smallest seq. #

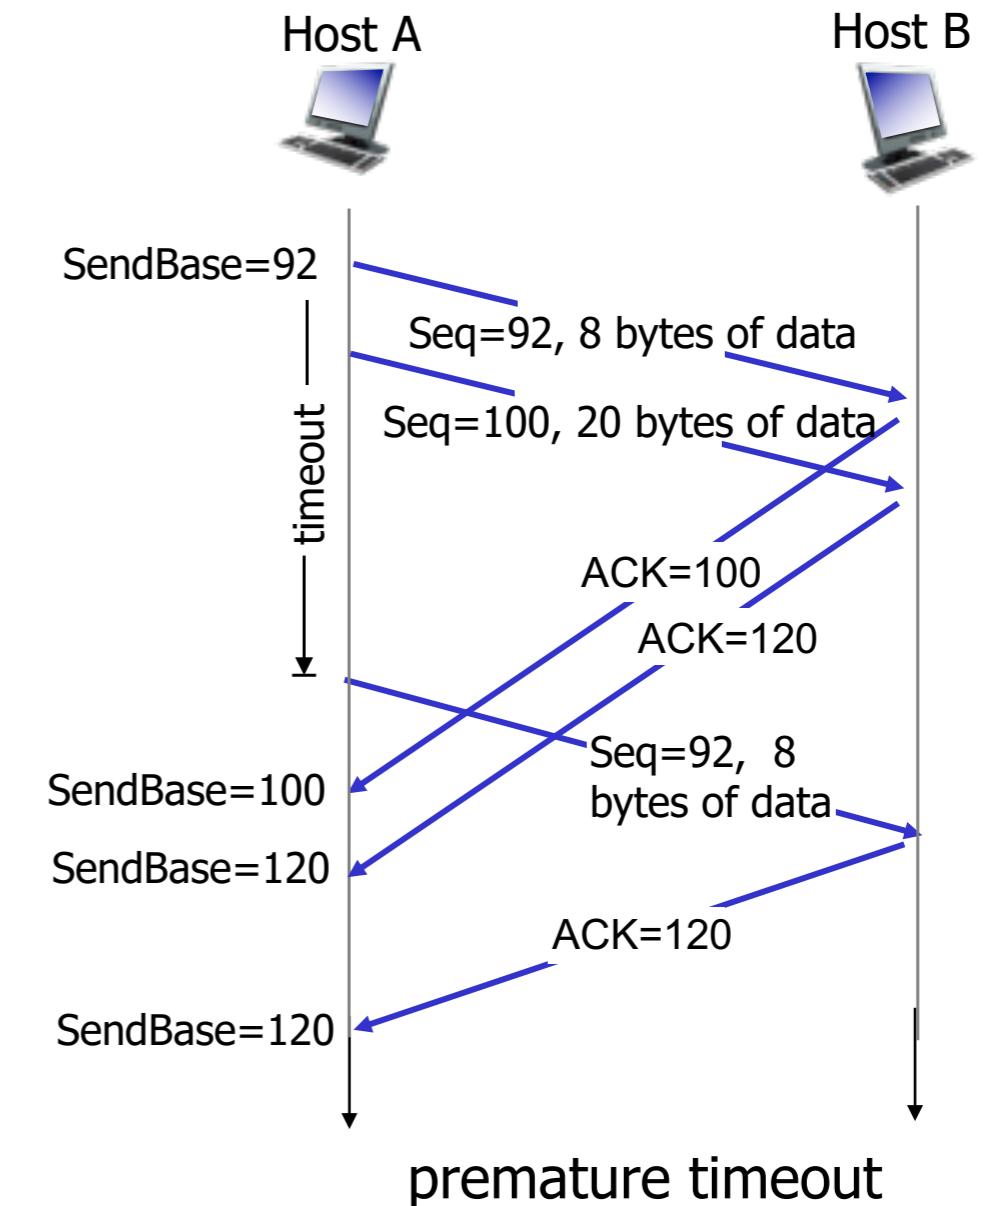
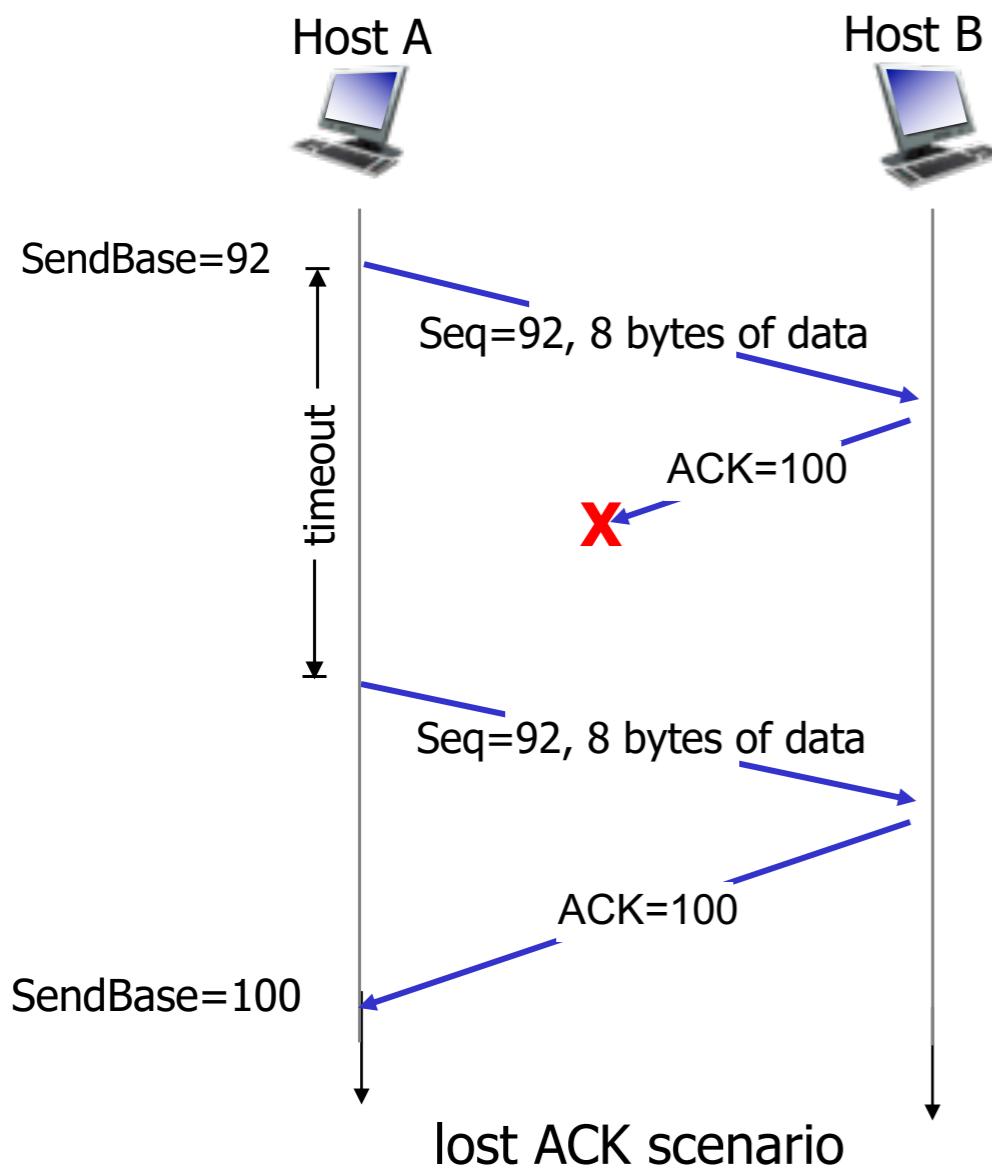
start timer

- retransmit oldest segment
that is not acked
- restart timer

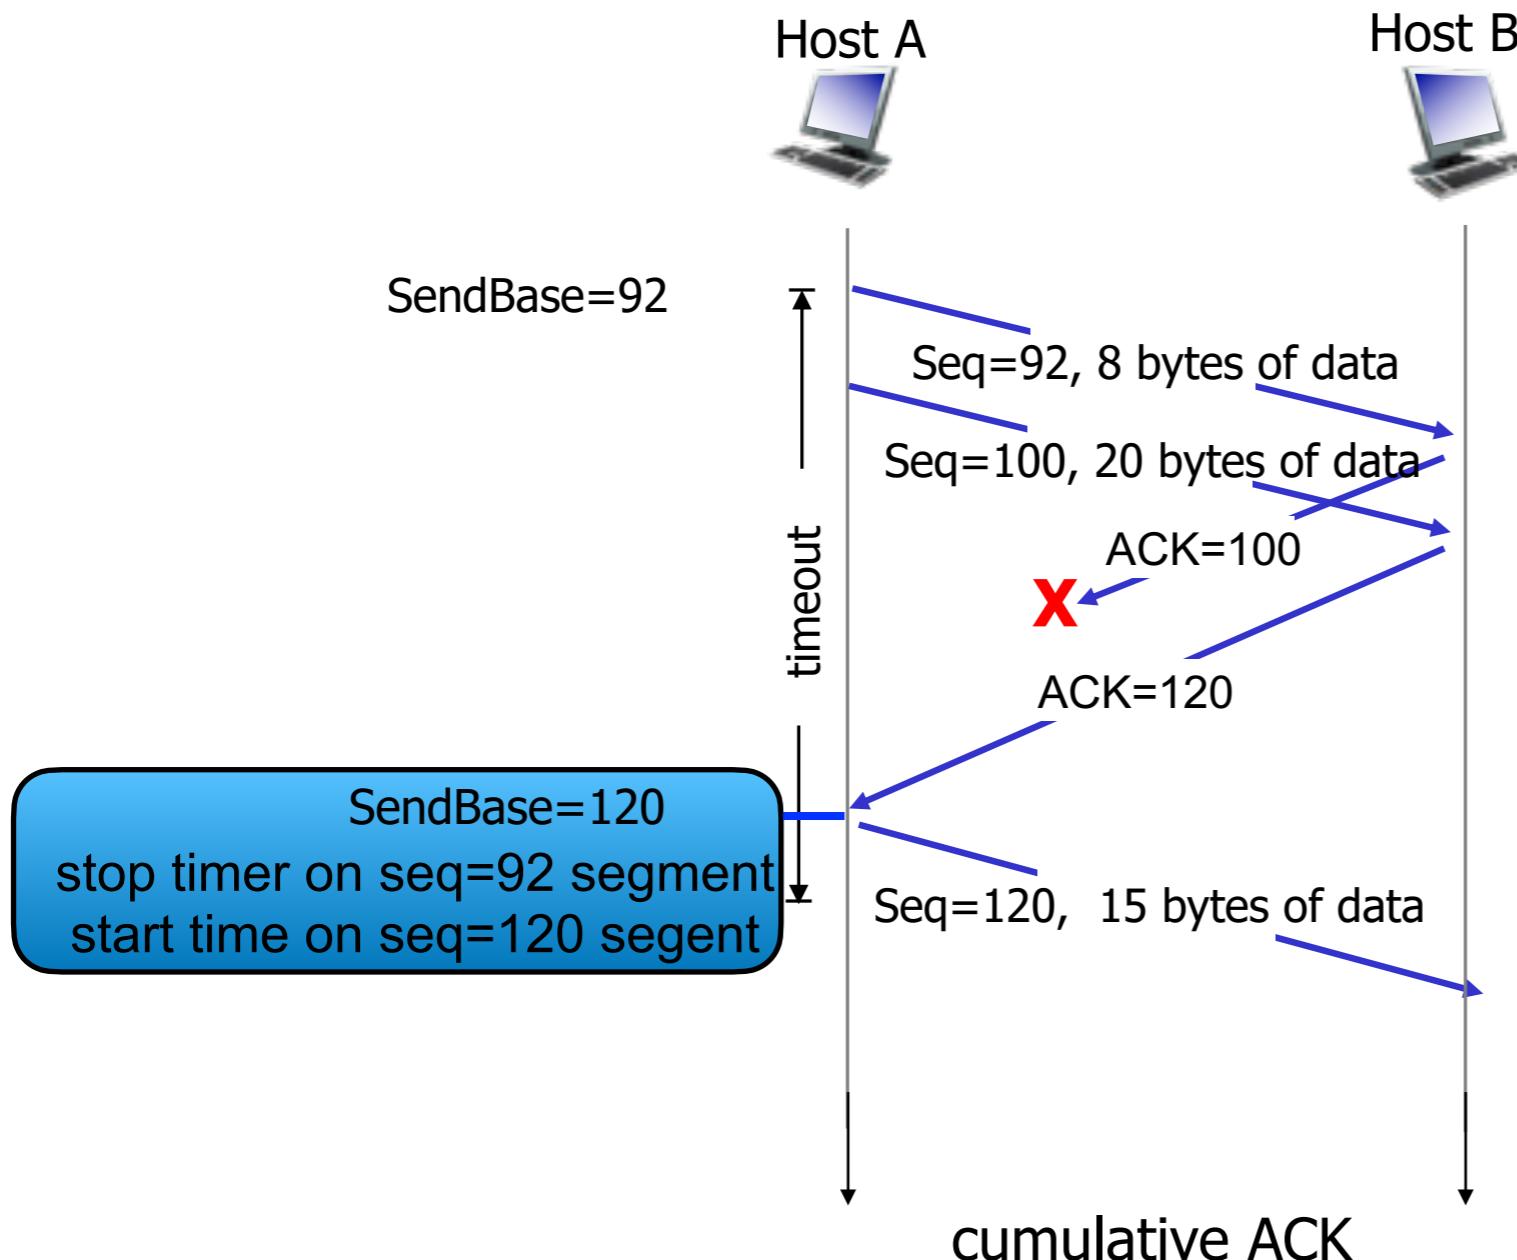
ack correctly received

- If acknowledging previously unacked
segments
advance send_base to next unacked seq. #
start timer if there are unacked segments

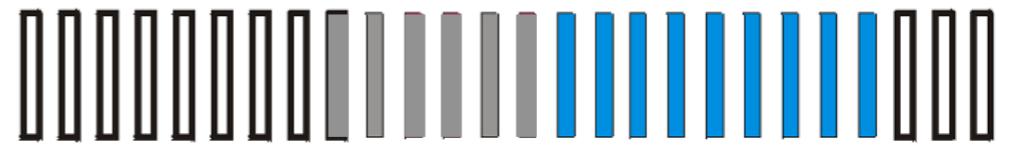
TCP RDT: retransmission scenarios



TCP RDT: retransmission scenarios



TCP RDT: receiver



TCP ACK Generation Recommendation [RFC 5681]

Event at Receiver	TCP Receiver action
Arrival of in-order segment with expected seq #. All data up to expected seq # already ACKed	Delayed ACK. Wait up to 500ms for another in-order segment. If no next segment, send an ACK.
Arrival of in-order segment with expected seq #. One other segment has ACK pending.	Immediately send single cumulative ACK, ACKing both in-order segments.
Arrival of out-of-order segment higher-than-expect seq. # . Gap detected.	Immediately send <i>duplicate ACK</i> , indicating seq. # of next expected byte.
Arrival of segment that partially or completely fills gap	Immediate send ACK, provided that segment starts at lower end of gap

TCP RDT: fast retransmit

- Detect lost segments via timeouts
 - timeout period often relatively long => long delay before resending lost packet
- Detect lost segments via **duplicate ACKs**
 - sender often sends many segments back-to-back
 - if segment is lost, there will likely be many duplicate ACKs.
- **TCP fast retransmit**: if sender receives 3 duplicate ACKs for the same data, it supposes that segment after ACKed data was lost, and resends segment before timeout

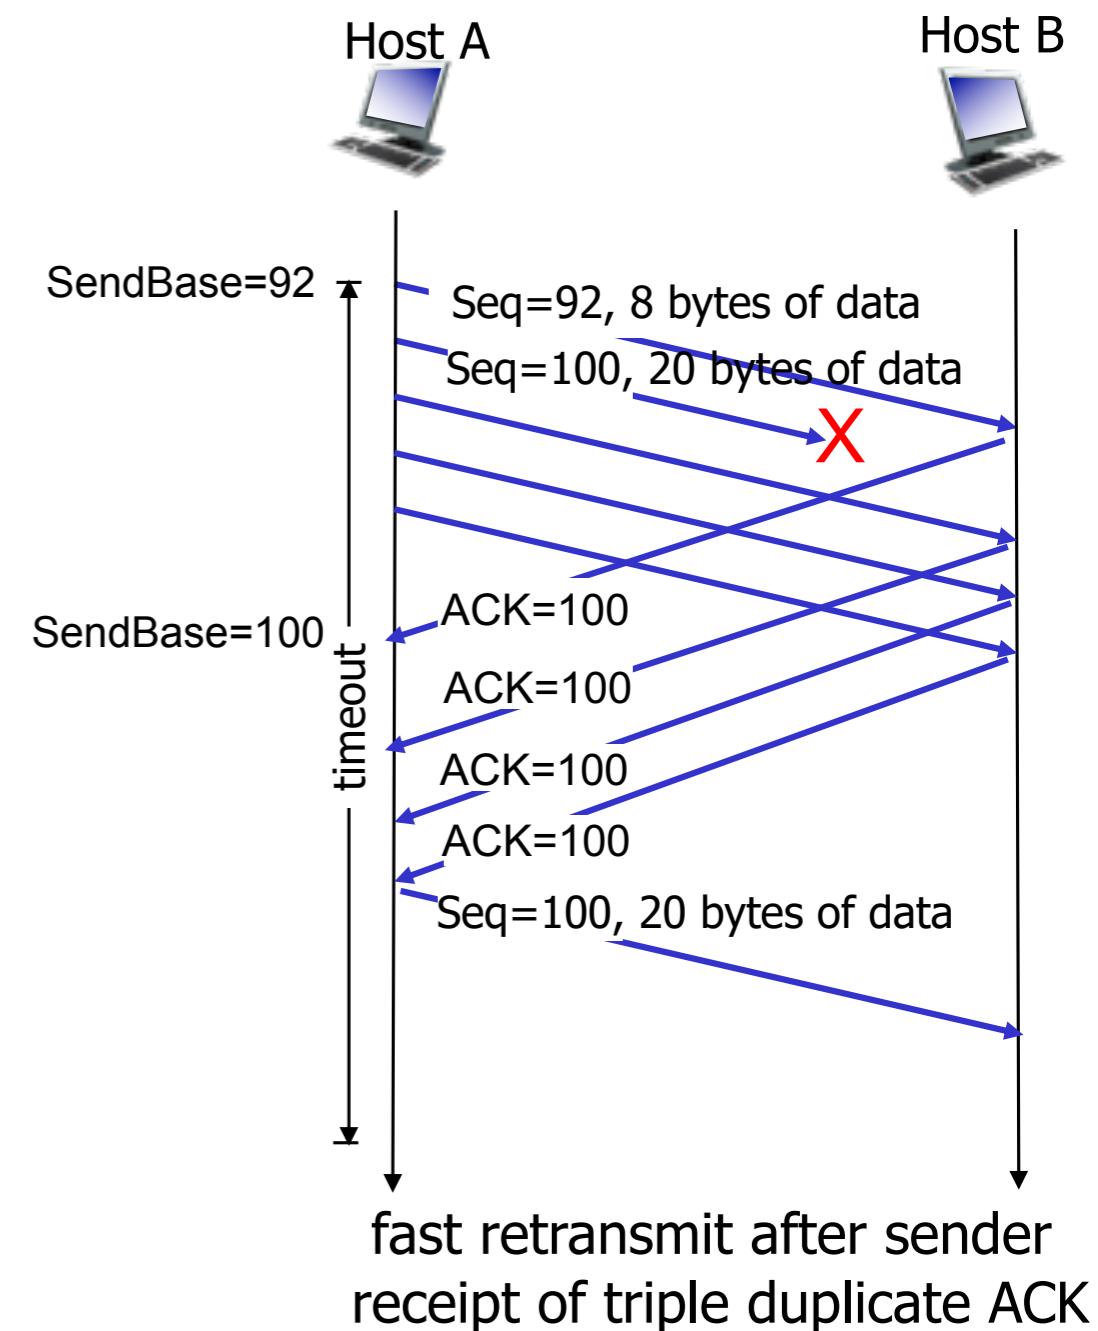
TCP RDT: fast retransmit (cont'd)

TCP sender with fast retransmit:
the following code replaces ACK received event on slide 9

```
event: ACK received, with ACK field value of y
    if (y > SendBase) {
        SendBase = y
        if (there are currently not-yet-acknowledged segments)
            start timer
        else stop time
    } else { //y=SendBase
        increment count of dup ACKs received for y
        if (count of dup ACKs received for y = 3) {
            resend segment with sequence number y
        }
    }
```

a duplicate ACK for
already ACKed segment

fast retransmit



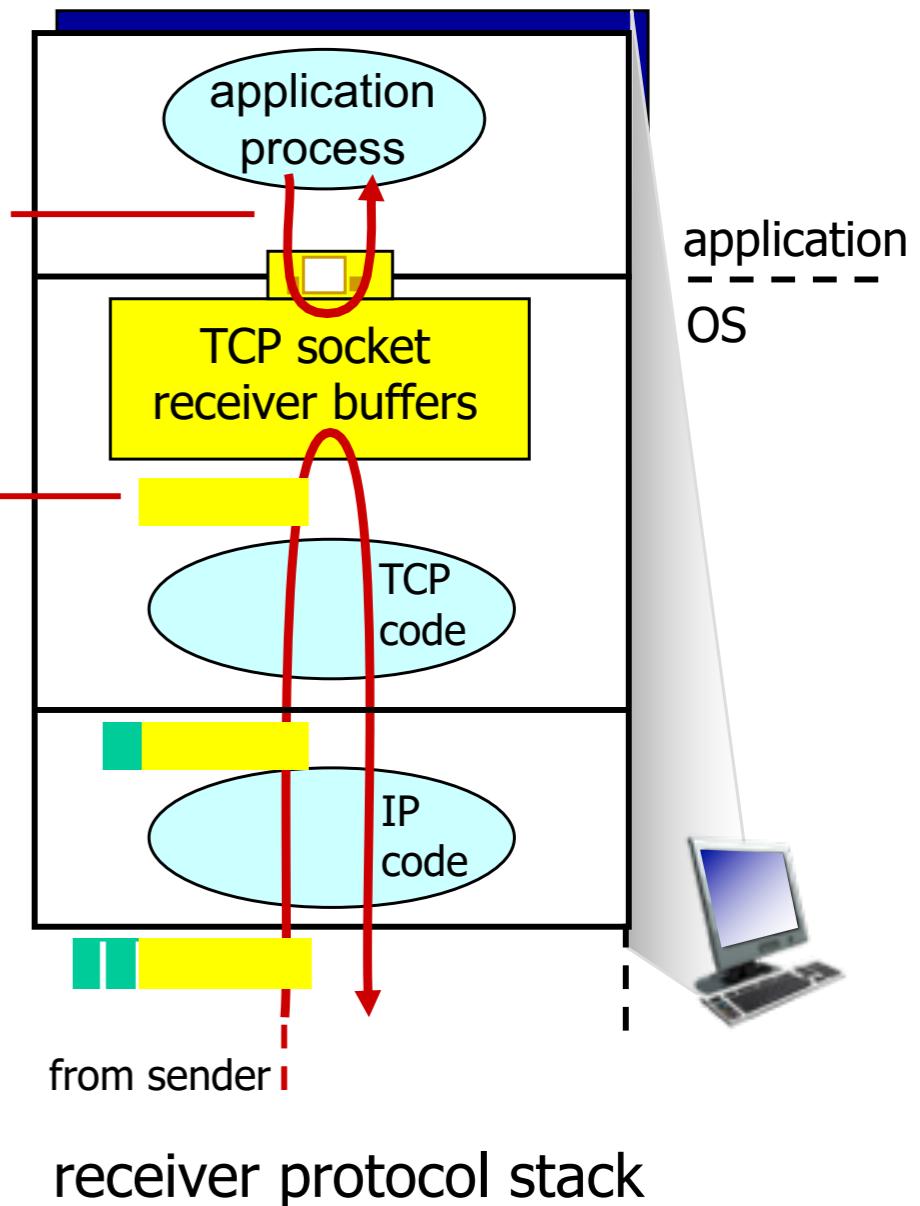
TCP flow control

- Receiver side of TCP connection has a receive buffer:

- Application process may be slow at reading from buffer
- **Flow control**
 - sender won't overflow receiver's buffer by sending too much, too fast
 - speed-matching service matching the sending rate to the receiving app's drain rate

application may remove data from TCP socket buffers

... may be slower than TCP
receiver is delivering (sender is sending)



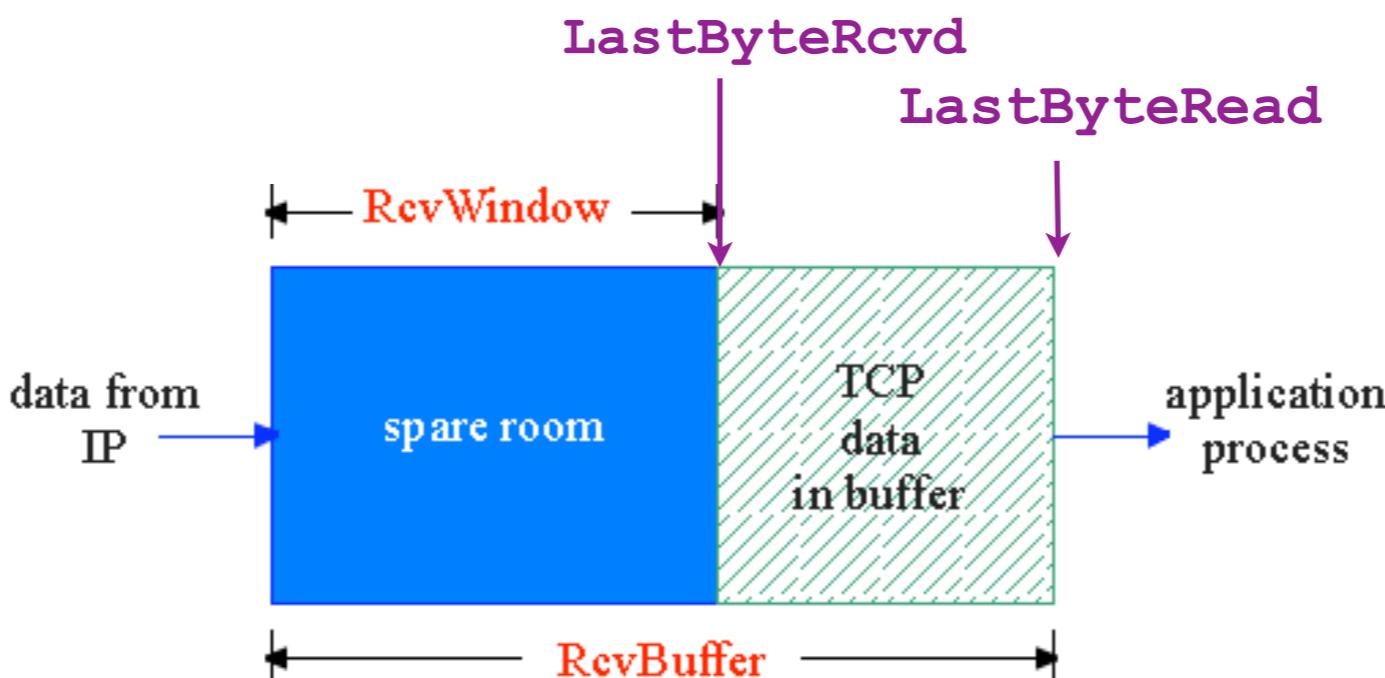
TCP flow control (cont'd)

□ How TCP flow control works

■ At TCP receiver

spare room in buffer = **RcvWindow** = **RcvBuffer**-
[**LastByteRcvd** - **LastByteRead**]

- Receiver advertises spare room by including value of current **RcvWindow** in segments to sender
RcvBuffer size can be set via socket options
- Sender limits unACKed data to **RcvWindow**
guarantees receive buffer does not overflow



TCP connection management: three-way handshake

Recall: TCP sender, receiver establish “connection” before exchanging data segments

How?

client: connection initiator
**clientSocket.connect(
 (serverName, serverPort))**
server: contacted by client
**connectionSocket, addr =
 serverSocket.accept()**

Why?

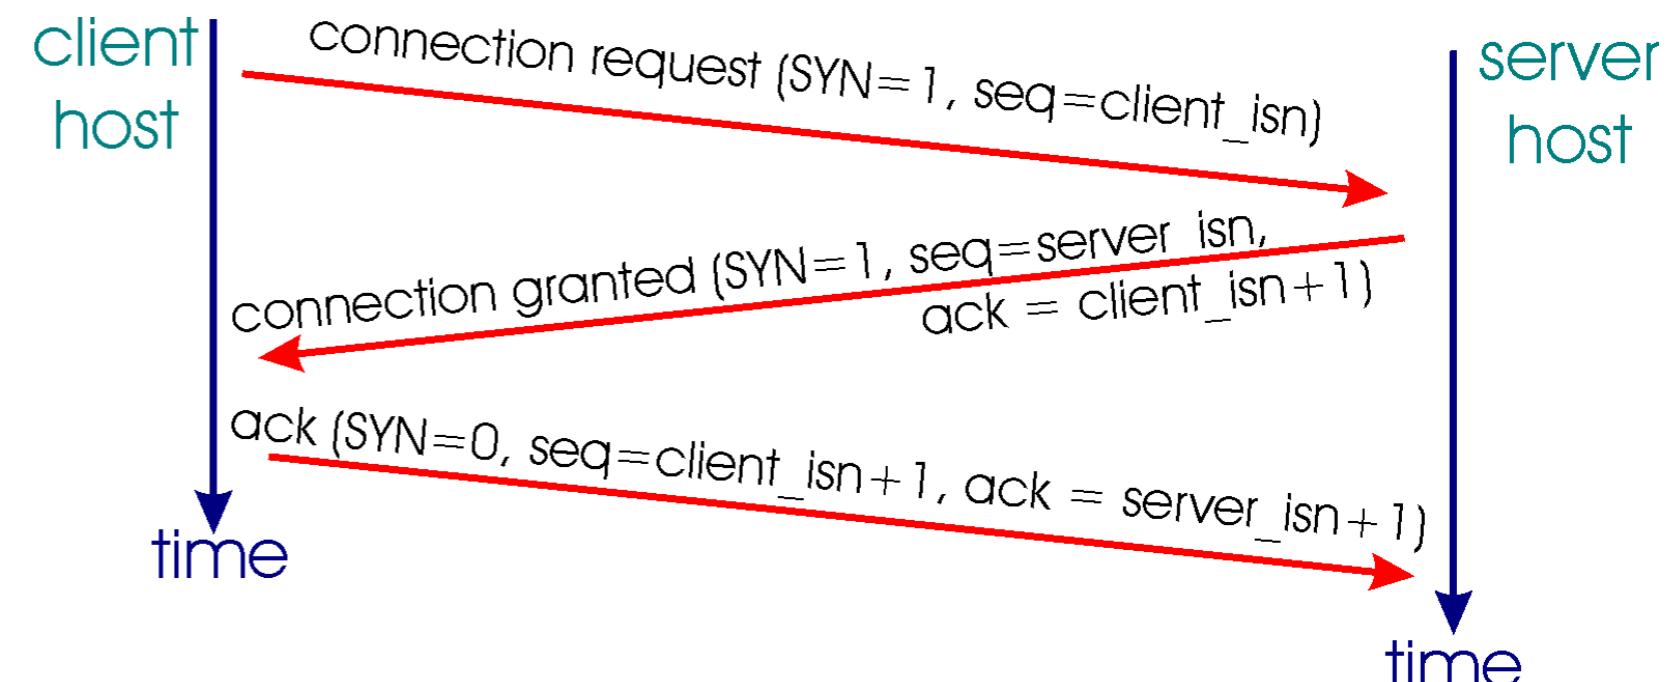
initialize TCP variables:
seq. #s
**RcvBuffer size at client,
 server**

Three-way handshake

Step 1: client host sends TCP SYN segment to server
★ specifies client initial seq #
★ no data

Step 2: server host receives SYN, replies with SYNACK segment
★ server allocates buffers
★ specifies server initial seq. #

Step 3: client receives SYNACK, replies with ACK segment, which may contain data



TCP connection management: closing a connection

Closing a connection (`socket.close()`)

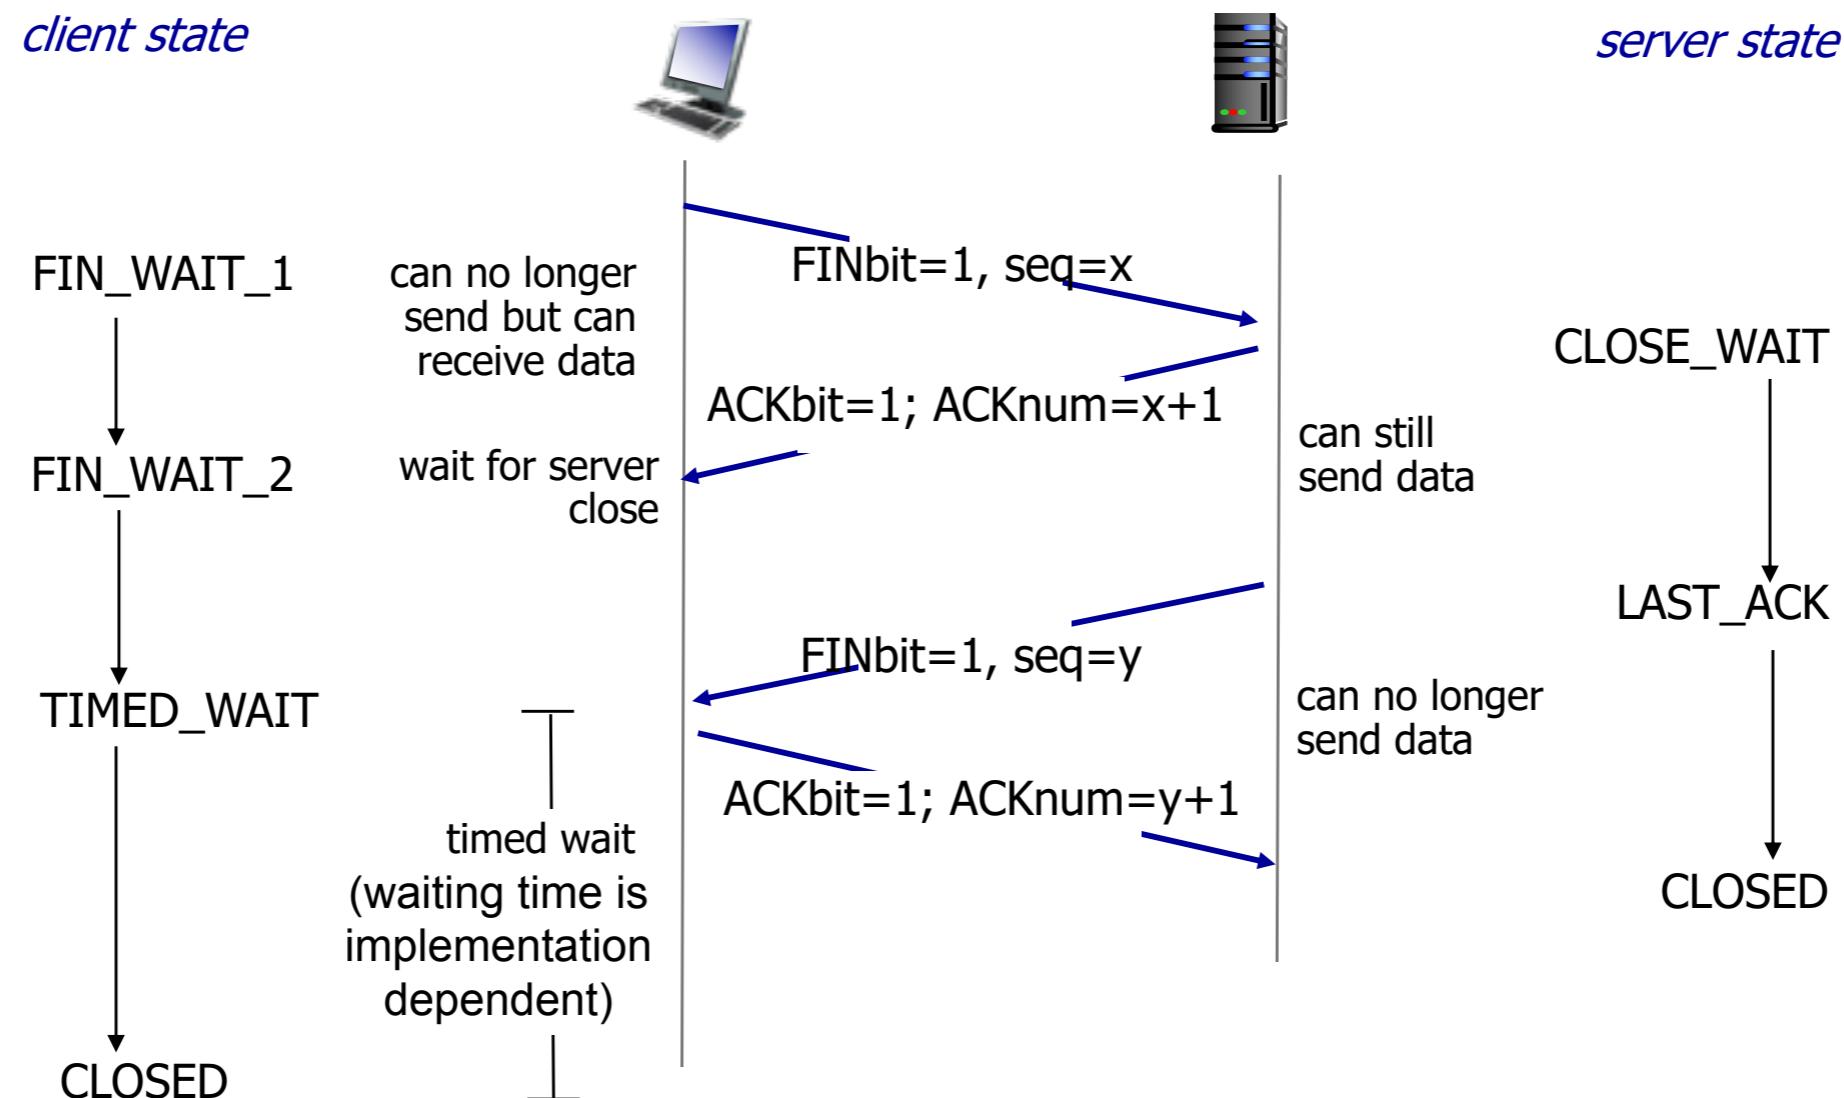
Step 1: client sends TCP FIN control segment (FIN bit is 1) to server

Step 2: server receives FIN, replies with ACK. Closing connection, sends FIN.

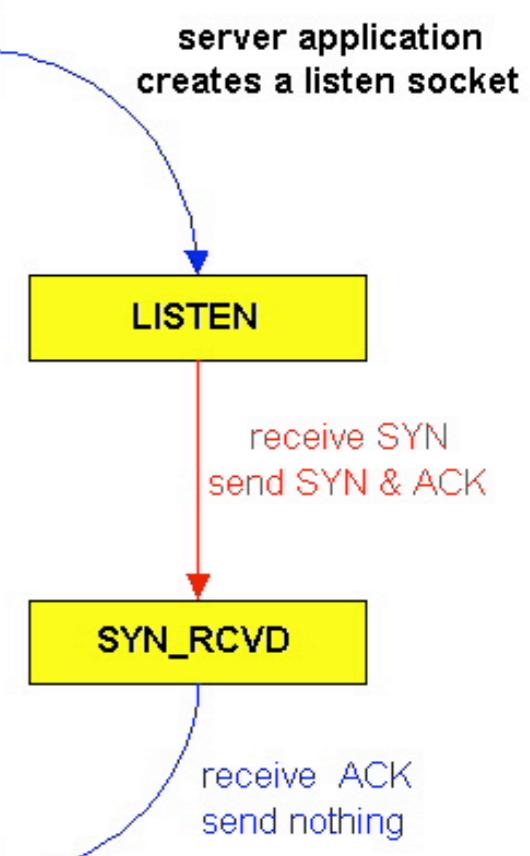
Step 3: client receives FIN, replies with ACK.

Enters “timed wait” - will respond with ACK to received FINs

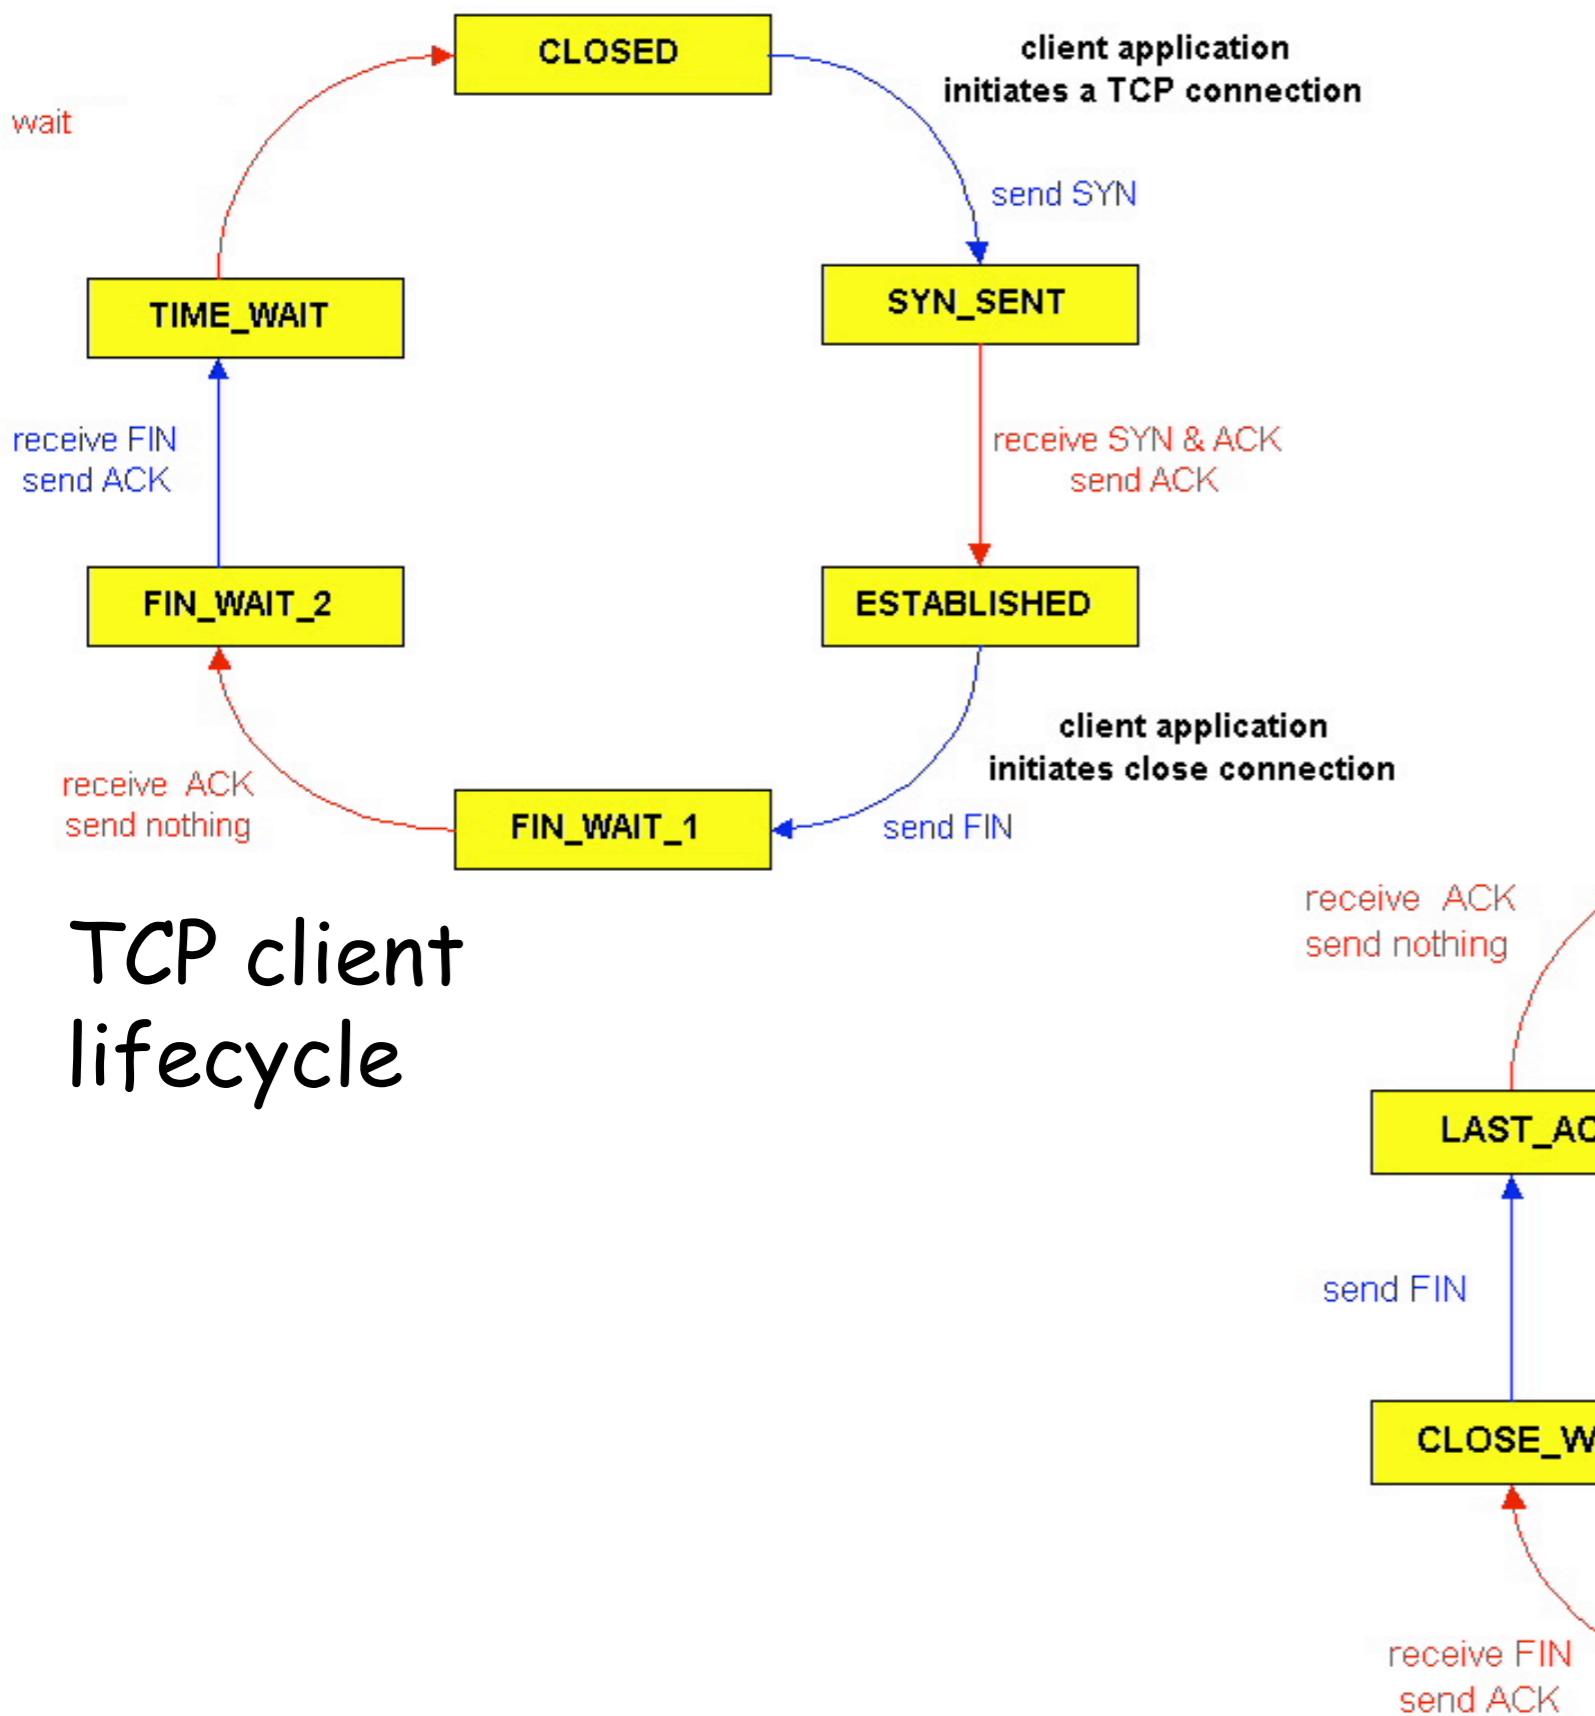
Step 4: server, receives ACK. Connection closed.



TCP server lifecycle



TCP client lifecycle



□ Required Reading:

- Computer Networking: A Top-Down Approach (8th Edition)
Ch 3.5.1, 3.5.2, 3.5.4, 3.5.5, 3.5.6

□ Acknowledgement:

- Some materials are extracted from the slides created by Prof. Jim F. Kurose and Prof. Keith W. Ross for the textbook.