



COMP 3234B

Computer and Communication Networks

2nd semester 2023-2024
Network Layer (III)


Prof. C Wu

Department of Computer Science
The University of Hong Kong

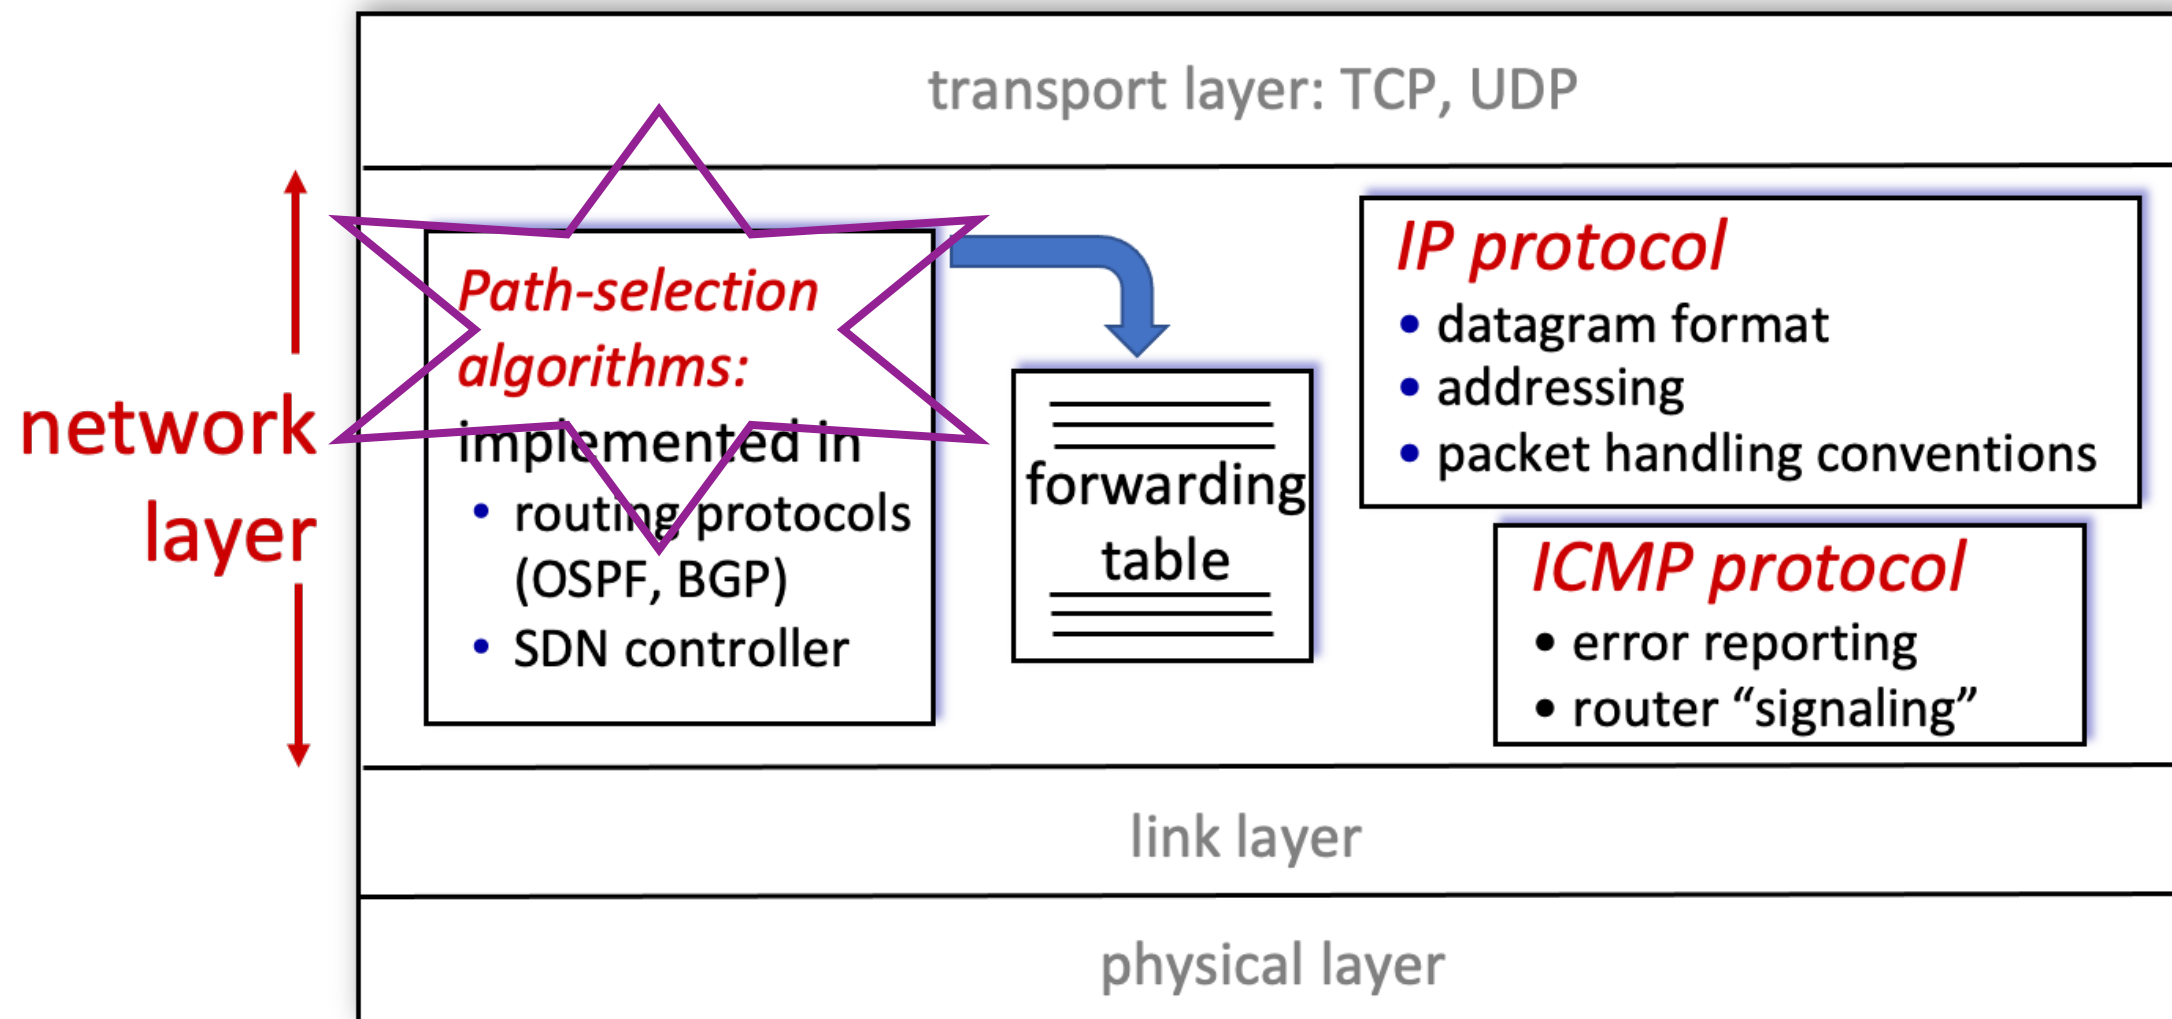
Roadmap

Network layer

- Principles behind network-layer services (ILO1)
forwarding vs. routing
network service models
- Router (ILO1)
- IP (ILO2,5)
DHCP
NAT
- ICMP(ILO2)
- **Routing algorithms (ILO3)**
- Routing in the Internet (ILO2,3)




Internet network layer



Routing protocols

- ❑ Decide the **good** paths (aka routes) taken by datagrams from sending host to receiving host, through the network of routers
 - decide the forwarding tables at routers
 - **path**: sequence of routers that datagrams traverse, going from source host to destination host
 - **“good”**: “least cost”, “fastest”, “least congested”



Graph abstraction of the network

□ Graph abstraction

■ Undirected graph: $G = (N, E)$

N = set of nodes (routers) = $\{ u, v, w, x, y, z \}$

E = set of edges (links) = $\{ (u,v), (u,x), (v,x), (v,w), (x,w), (x,y), (w,y), (w,z), (y,z) \}$


■ Link cost: the value associated with a link (related to bandwidth, congestion, physical length, delay, etc.)

$c(x_1, x_2)$ = cost of link (x_1, x_2) , e.g., $c(w, z) = 5$

$c(x_1, x_2) = \infty$ if (x_1, x_2) does not belong to E

■ Path: a sequence of nodes, $(x_1, x_2, x_3, \dots, x_p)$

■ Cost of path $(x_1, x_2, x_3, \dots, x_p)$: $c(x_1, x_2) + c(x_2, x_3) + \dots + c(x_{p-1}, x_p)$



What's the least-cost path between u and z ?

□ Routing algorithm

■ Given a set of routers with links connecting the routers, finds a least-cost path from source router to destination router

(finds the **shortest path** when link costs represent length)

Routing algorithm classification

□ Centralized or decentralized routing algorithm

■ Centralized:

algorithm input: complete global knowledge about the network, including node connectivity (topology), link costs

carried out at each router


link-state (LS) algorithms

■ Decentralized:

with (local) knowledge of connectivity to neighbor routers, link costs to neighbors

algorithm carried out at each router by **iterative** process of computation and exchange of information with neighbor


distance-vector (DV) algorithms



Link-state routing algorithm

□ Dijkstra's algorithm

- Input: network topology, link costs of all links
- Output: least-cost paths from **one** node (“source router”) to all other nodes (routers)
derives forwarding table for that router
- **Iterative algorithm executed at one node:**
after k iterations, the source derives
least-cost paths to k destinations with
the smallest path costs



each node broadcasts its identity and costs of neighboring links to all other nodes in the network (“link state broadcast”)
=>all nodes have same and complete view of the network

Dijkstra's algorithm

□ Notation

- $c(x,y)$: link cost from node x to y ;
 $= \infty$ if not direct neighbors
- $D(v)$: current value of cost of path from source to destination v
- $p(v)$: predecessor node along the current least-cost path from source to v
- N' : set of nodes whose least-cost path already known


1 **Initialization:**

```
2   $N' = \{u\}$ 
3  for all nodes  $i$ 
4    if  $i$  is a neighbor of  $u$ 
5      then  $D(i) = c(u,i)$ ,  $p(i) = u$ 
6    else  $D(i) = \infty$ 
```

8 **Loop**


```
9  find  $j$  not in  $N'$  such that  $D(j)$  is minimum
10 add  $j$  to  $N'$ 
11 update  $D(i)$  for each neighbor  $i$  of  $j$  and not in  $N'$  :
12    $D(i) = \min( D(i), D(j) + c(j,i) );$  update  $p(i)$ 
13   /* new cost to  $i$  is either old cost to  $i$  or known
14   least path cost to  $j$  plus cost from  $j$  to  $i$  */
15 until  $N' = N$ 
```

of loops=
of nodes in the network
excluding the source



An example of Dijkstra's algorithm


□ Example



Step	N'	D(v),p(v)	D(w),p(w)	D(x),p(x)	D(y),p(y)	D(z),p(z)
0	u	2,u	5,u	1,u	∞	∞
1	ux	2,u	4,x		2,x	∞
2	uxy	2,u	3,y			4,y
3	uxyv		3,y			4,y
4	uxyvw					4,y
5	uxyvwz					

An example of Dijkstra's algorithm (cont'd)

Resulting least-cost-path tree from u:



Resulting forwarding table in u:

destination	link
v	(u,v)
x	(u,x)
y	(u,x)
w	(u,x)
z	(u,x)

next-hop node along the
least-cost path towards
the destination

Discussions on Dijkstra's algorithm

□ Algorithm complexity

with n nodes (routers excluding the source)

■ each iteration: need to check all nodes, w , not in N'

■ $n(n+1)/2$ comparisons: $O(n^2)$

■ more efficient implementations possible: $O(n \log n)$

Discussions on Dijkstra's algorithm (cont'd)

□ Potential problem: **routing oscillations**

■ Example scenario:

link cost = amount of carried traffic (reflecting delay on the link)


link costs could be asymmetric: $c(u,v)=c(v,u)$ only if traffic on both directions is the same

node D originates a unit of traffic destined to A;


node B originates a unit of traffic destined to A;

node C originates traffic of amount e to A.


■ Possible solution: have each router run the algorithm at different times




initially
(new costs shown after
initial routing)



given these costs,
find new routing....
resulting in new costs



given these costs,
find new routing....
resulting in new costs



given these costs,
find new routing....
resulting in new costs

Distance-vector routing algorithm

□ Bellman-Ford algorithm


- Input: connectivity/link costs to neighbors
- Output: least-cost paths from each node to all other nodes

derives forwarding table for each router

- Iterative, asynchronous, distributed algorithm executed by all nodes together:

each node receives updates from neighbors, recomputes, and distributes its new calculation result to neighbors

algorithm terminates (least-cost path from each node to each other node derived) when no more update is exchanged between neighbors



each node updates
information to
neighbors only

Bellman-Ford equation

□ Bellman-Ford equation


- Define an important relationship among the costs of least-cost paths

$d_x(y) :=$ cost of least-cost path from x to y

Then


$$d_x(y) = \min_v \{c(x,v) + d_v(y)\}$$

where min is taken over all neighbors v of x



Bellman-Ford equation (cont'd)

■ Example



$$d_v(z) = 5, d_x(z) = 3, d_w(z) = 3$$

Bellman-Ford equation says:

$$d_u(z) = \min \{ c(u,v) + d_v(z), \\ c(u,x) + d_x(z), \\ c(u,w) + d_w(z) \}$$

$$= \min \{ 2 + 5, \\ 1 + 3, \\ 5 + 3 \} = 4$$

Node that achieves minimum is next hop in the least-cost path
→ decides the entry in u's forwarding table

Bellman-Ford algorithm basics

□ $D_x(y)$ = estimate of least path cost from x to y

□ Node x maintains

- cost to each neighbor v: $c(x,v)$
- distance vector $D_x = [D_x(y): y \in N]$
- its neighbors' distance vectors

For each neighbor v, x maintains $D_v = [D_v(y): y \in N]$

□ Basic idea of Bellman-Ford algorithm

- From time-to-time, each node sends its updated distance vector (DV) to neighbors
- When a node x receives new DV estimate from neighbor v, it updates its own DV using Bellman-Ford equation:

$$D_x(y) \leftarrow \min_v \{c(x,v) + D_v(y)\} \quad \text{for each node } y \in N$$

- all nodes continue to exchange their DVs in an asynchronous fashion; each least cost estimate $D_x(y)$ converges to the actual least cost $d_x(y)$

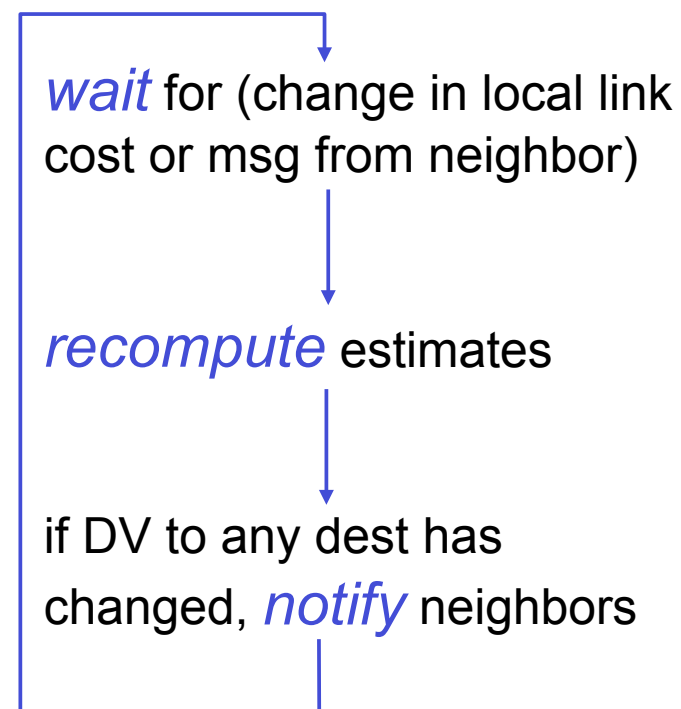
Bellman-Ford algorithm

At each node, x:

1 **Initialization:**

- 2 for all destinations y in N:
- 3 $D_x(y) = c(x,y)$
- 4 for each neighbor w
- 5 $D_w(y) = \infty$ for all destinations y in N
- 6 for each neighbor w
- 7 send distance vector $\mathbf{D}_x = [D_x(y): y \in N]$ to w
- 8
- 9 **Loop**
- 10 wait (until I see a link cost change to some neighbor w or
- 11 until I receive a distance vector from some neighbor w)
- 12
- 13 for each y in N:
- 14 $D_x(y) = \min_v \{c(x,v) + D_v(y)\}$
- 15
- 16 if $D_x(y)$ changed for any destination y
- 17 send distance vector $\mathbf{D}_x = [D_x(y): y \in N]$ to all neighbors
- 18
- 19 **forever**

Each node:



An example of Bellman-Ford algorithm

□ Example

$$D_x(y) = \min\{c(x,y) + D_y(y), c(x,z) + D_z(y)\}$$

$$= \min\{2+0, 7+1\} = 2$$

$$D_x(z) = \min\{c(x,y) + D_y(z), c(x,z) + D_z(z)\}$$

$$= \min\{2+1, 7+0\} = 3$$

node x table

		cost to		
		x	y	z
from	x	0	2	7
	y	∞	∞	∞
	z	∞	∞	∞

node y table

		cost to		
		x	y	z
from	x	∞	∞	∞
	y	2	0	1
	z	∞	∞	∞


node z table

		cost to		
		x	y	z
from	x	∞	∞	∞
	y	∞	∞	∞
	z	7	1	0

		cost to		
		x	y	z
from	x	0	2	3
	y	2	0	1
	z	7	1	0

		cost to		
		x	y	z
from	x	0	2	7
	y	2	0	1
	z	7	1	0

		cost to		
		x	y	z
from	x	0	2	7
	y	2	0	1
	z	3	1	0



time

An example of Bellman-Ford algorithm (cont'd)

node x table

		cost to		
		x	y	z
from	x	0	2	7
	y	∞	∞	∞
	z	∞	∞	∞

node y table

		cost to		
		x	y	z
from	x	∞	∞	∞
	y	2	0	1
	z	∞	∞	∞

node z table

		cost to		
		x	y	z
from	x	∞	∞	∞
	y	∞	∞	∞
	z	7	1	0

		cost to		
		x	y	z
from	x	0	2	3
	y	2	0	1
	z	7	1	0


		cost to		
		x	y	z
from	x	0	2	7
	y	2	0	1
	z	7	1	0

		cost to		
		x	y	z
from	x	0	2	7
	y	2	0	1
	z	3	1	0

		cost to		
		x	y	z
from	x	0	2	3
	y	2	0	1
	z	3	1	0

		cost to		
		x	y	z
from	x	0	2	3
	y	2	0	1
	z	3	1	0

		cost to		
		x	y	z
from	x	0	2	3
	y	2	0	1
	z	3	1	0




time

Discussions on Bellman-Ford algorithm I

□ Link cost change: scenario 1

- At time t_0 , x and y detect the link-cost change ($4 \rightarrow 1$), updates their DV, and inform neighbors.



DV table evolution

node x table


		cost to		
		x	y	z
from	x	0	4 ¹	5 ²
	y	4	0	1
	z	5	1	0

node y table

		cost to		
		x	y	z
from	x	0	4	5
	y	4 ¹	0	1
	z	5	1	0

node z table

		cost to		
		x	y	z
from	x	0	4	5
	y	4	0	1
	z	5	1	0



t₀

t₁

t₂

time

DV table evolution

node x table

		cost to		
		x	y	z
from	x	0	4 ¹	5 ²
	y	4	0	1
	z	5	1	0

		cost to		
		x	y	z
from	x	0	1	2
	y	1	0	1
	z	5	1	0

node y table


		cost to		
		x	y	z
from	x	0	4	5
	y	4 ¹	0	1
	z	5	1	0

		cost to		
		x	y	z
from	x	0	1	2
	y	1	0	1
	z	5	1	0

node z table

		cost to		
		x	y	z
from	x	0	4	5
	y	4	0	1
	z	5	1	0

		cost to		
		x	y	z
from	x	0	1	2
	y	1	0	1
	z	2	1	0



t₀

t₁

t₂

time

DV table evolution

node x table

		cost to		
		x	y	z
from	x	0	4	5
	y	4	0	1
	z	5	1	0

		cost to		
		x	y	z
from	x	0	1	2
	y	1	0	1
	z	5	1	0

		cost to		
		x	y	z
from	x	0	1	2
	y	1	0	1
	z	2	1	0

node y table

		cost to		
		x	y	z
from	x	0	4	5
	y	4	0	1
	z	5	1	0

		cost to		
		x	y	z
from	x	0	1	2
	y	1	0	1
	z	5	1	0


		cost to		
		x	y	z
from	x	0	1	2
	y	1	0	1
	z	2	1	0

node z table

		cost to		
		x	y	z
from	x	0	4	5
	y	4	0	1
	z	5	1	0

		cost to		
		x	y	z
from	x	0	1	2
	y	1	0	1
	z	2	1	0

		cost to		
		x	y	z
from	x	0	1	2
	y	1	0	1
	z	2	1	0




t₀

t₁

t₂

time

Discussions on Bellman-Ford algorithm I (cont'd)



□ Link cost change: scenario 1


- At time t_0 , x and y detect the link-cost change ($4 \rightarrow 1$), updates their DV, and inform neighbors.
- At time t_1 , z receives the updates, computes a new least cost to x ($5 \rightarrow 2$) and informs its neighbors.
- At time t_2 , x and y receive z's update and update their DV tables. x and y's least costs do not change and hence do not send any further message to z.

Good news travels fast!

Discussions on Bellman-Ford algorithm II

□ Link cost change: scenario 2

- At time t_0 , x and y detect the link-cost change ($4 \rightarrow 60$), update their DV, and notify neighbors




$$D_x(y) = \min\{c(x,y) + D_y(y), D_y(x) = \min\{c(y,x) + D_x(x), c(x,z) + D_z(y)\}$$

$$= \min\{60+0, 50+1\} = 51$$

$$D_y(x) = \min\{c(y,x) + D_x(x), c(y,z) + D_z(x)\}$$

$$= \min\{60+0, 1+5\} = 6$$



node x table

		cost to		
		x	y	z
from	x	0	4	5
	y	4	0	1
	z	5	1	0

node y table

		cost to		
		x	y	z
from	x	0	4	5
	y	4	0	1
	z	5	1	0

node z table

		cost to		
		x	y	z
from	x	0	4	5
	y	4	0	1
	z	5	1	0

t₀

t₁

t₂


t₃

t₄

time

$$D_z(x) = \min\{c(z,x) + D_x(x), c(z,y) + D_y(x)\}$$

$$= \min\{50 + 0, 1 + 6\} = 7$$



node x table

		cost to		
		x	y	z
from	x	0	4	5
	y	4	0	1
	z	5	1	0

		cost to		
		x	y	z
from	x	0	51	50
	y	6	0	1
	z	5	1	0

node y table

		cost to		
		x	y	z
from	x	0	4	5
	y	4	0	1
	z	5	1	0

		cost to		
		x	y	z
from	x	0	51	50
	y	6	0	1
	z	5	1	0

node z table

		cost to		
		x	y	z
from	x	0	4	5
	y	4	0	1
	z	5	1	0

		cost to		
		x	y	z
from	x	0	51	50
	y	6	0	1
	z	7	1	0

t₀

t₁

t₂


t₃

t₄

time

$$D_y(x) = \min\{c(y,x) + D_x(x), c(y,z) + D_z(x)\}$$

$$= \min\{60 + 0, 1 + 7\} = 8$$



node x table

		cost to		
		x	y	z
from	x	0	4	5
	y	4	0	1
	z	5	1	0

		cost to		
		x	y	z
from	x	0	51	50
	y	6	0	1
	z	5	1	0

		cost to		
		x	y	z
from	x	0	51	50
	y	6	0	1
	z	7	1	0

node y table

		cost to		
		x	y	z
from	x	0	4	5
	y	4	0	1
	z	5	1	0

		cost to		
		x	y	z
from	x	0	51	50
	y	6	0	1
	z	5	1	0

		cost to		
		x	y	z
from	x	0	51	50
	y	8	0	1
	z	7	1	0

node z table

		cost to		
		x	y	z
from	x	0	4	5
	y	4	0	1
	z	5	1	0


		cost to		
		x	y	z
from	x	0	51	50
	y	6	0	1
	z	7	1	0

		cost to		
		x	y	z
from	x	0	51	50
	y	6	0	1
	z	7	1	0

t_0
 t_1
 t_2
 t_3
 t_4
time

$$D_z(x) = \min\{c(z,x) + D_x(x), c(z,y) + D_y(x)\}$$

$$= \min\{50 + 0, 1 + 8\} = 9$$



node x table

		cost to		
		x	y	z
from	x	0	4	5
	y	4	0	1
	z	5	1	0

		cost to		
		x	y	z
from	x	0	51	50
	y	6	0	1
	z	5	1	0

		cost to		
		x	y	z
from	x	0	51	50
	y	6	0	1
	z	7	1	0

		cost to		
		x	y	z
from	x	0	51	50
	y	8	0	1
	z	7	1	0

node y table

		cost to		
		x	y	z
from	x	0	4	5
	y	4	0	1
	z	5	1	0

		cost to		
		x	y	z
from	x	0	51	50
	y	6	0	1
	z	5	1	0

		cost to		
		x	y	z
from	x	0	51	50
	y	8	0	1
	z	7	1	0

		cost to		
		x	y	z
from	x	0	51	50
	y	8	0	1
	z	7	1	0

node z table

		cost to		
		x	y	z
from	x	0	4	5
	y	4	0	1
	z	5	1	0

		cost to		
		x	y	z
from	x	0	51	50
	y	6	0	1
	z	7	1	0

		cost to		
		x	y	z
from	x	0	51	50
	y	6	0	1
	z	7	1	0

		cost to		
		x	y	z
from	x	0	51	50
	y	8	0	1
	z	9	1	0

t₀

t₁

t₂


t₃

t₄

time

$$D_y(x) = \min\{c(y,x) + D_x(x), c(y,z) + D_z(x)\}$$

$$= \min\{60 + 0, 1 + 9\} = 10$$



node x table

		cost to		
		x	y	z
from	x	0	4	5
	y	4	0	1
	z	5	1	0

		cost to		
		x	y	z
from	x	0	51	50
	y	6	0	1
	z	5	1	0

		cost to		
		x	y	z
from	x	0	51	50
	y	6	0	1
	z	7	1	0

		cost to		
		x	y	z
from	x	0	51	50
	y	8	0	1
	z	7	1	0

		cost to		
		x	y	z
from	x	0	51	50
	y	8	0	1
	z	9	1	0

node y table

		cost to		
		x	y	z
from	x	0	4	5
	y	4	0	1
	z	5	1	0

		cost to		
		x	y	z
from	x	0	51	50
	y	6	0	1
	z	5	1	0

		cost to		
		x	y	z
from	x	0	51	50
	y	8	0	1
	z	7	1	0

		cost to		
		x	y	z
from	x	0	51	50
	y	8	0	1
	z	7	1	0

		cost to		
		x	y	z
from	x	0	51	50
	y	10	0	1
	z	9	1	0

....

node z table


		cost to		
		x	y	z
from	x	0	4	5
	y	4	0	1
	z	5	1	0

		cost to		
		x	y	z
from	x	0	51	50
	y	6	0	1
	z	7	1	0

		cost to		
		x	y	z
from	x	0	51	50
	y	6	0	1
	z	7	1	0

		cost to		
		x	y	z
from	x	0	51	50
	y	8	0	1
	z	9	1	0


		cost to		
		x	y	z
from	x	0	51	50
	y	8	0	1
	z	9	1	0



Discussions on Bellman-Ford algorithm II (cont'd)

❑ Link cost change: scenario 2

- At time t_0 , x and y detect the link-cost change ($4 \rightarrow 60$), update their DV, and notify neighbors
- At time t_1 , z receives updated DVs, computes a new least cost to x of $D_z(x) = \min \{50+0, 1+6\} = 7$, and informs neighbors of its new DV
- At time t_2 , y receives z's update, recomputes $D_y(x) = 8$, and sends it to neighbors
- At time t_3 , z receives y's update, recomputes $D_z(x) = 9$, and sends it to neighbors
- ...




44 iterations before algorithm terminates!

Bad news travels slowly => "count to infinity" problem!

But it does not solve general count-to-infinity problem

The problem in this scenario can be avoided by **poisoned reverse**:
If Z routes through Y to get to X, Z tells Y its (Z's) distance to X is infinite (so Y won't route to X via Z)

Poisoned reverse



node x table

		cost to					cost to					cost to					cost to					cost to					cost to		
		x	y	z			x	y	z			x	y	z			x	y	z			x	y	z			x	y	z
from	x	0	4	5	from	x	0	51	50	from	x	0	51	50	from	x	0	51	50	from	x	0	51	50	from	x	0	51	50
	y	4	0	1		y	6	0	1		y	6	0	1		y	60	0	1		y	60	0	1		y	51	0	1
	z	5	1	0		z	5	1	0		z	7	1	0		z	7	1	0		z	50	1	0		z	50	1	0

node y table

		cost to					cost to					cost to					cost to					cost to					cost to		
		x	y	z			x	y	z			x	y	z			x	y	z			x	y	z			x	y	z
from	x	0	4	5	from	x	0	51	50	from	x	0	51	50	from	x	0	51	50	from	x	0	51	50	from	x	0	51	50
	y	4	0	1		y	6	0	1		y	60	0	1		y	51	0	1		y	51	0	1		y	51	0	1
	z	5	1	0		z	5	1	0		z	∞	1	0		z	50	1	0		z	50	1	0		z	50	1	0

node z table

		cost to					cost to					cost to					cost to					cost to					cost to		
		x	y	z			x	y	z			x	y	z			x	y	z			x	y	z			x	y	z
from	x	0	4	5	from	x	0	51	50	from	x	0	51	50	from	x	0	51	50	from	x	0	51	50	from	x	0	51	50
	y	4	0	1		y	6	0	1		y	6	0	1		y	60	0	1		y	60	0	1		y	51	0	1
	z	5	1	0		z	7	1	0		z	7	1	0		z	50	1	0		z	50	1	0		z	50	1	0

t_0

t_1

t_2

t_3

t_4

t_5

z routes to x through y; suppose it tells y that $D_z(x) = \infty$ in next round

Comparison of LS and DV algorithms

Message complexity

LS: with n nodes, $O(n^2)$ msgs

DV: exchange between neighbors only;
convergence time varies

Convergence

LS: $O(n^2)$ algorithm requires $O(n^2)$ msgs;
may have oscillations

DV: convergence time varies;
may existing routing loops
(count-to-infinity problem)

Both algorithms used in routing protocols in the Internet

LS algorithms used in OSPF, IS-IS.

DV algorithms used in RIP, IGRP.

❑ Required reading:

- *Computer Networking: A Top Down Approach* (8th Edition)
Ch 5.2

❑ Acknowledgement:

- Some materials are extracted from the slides created by Prof. Jim F. Kurose and Prof. Keith W. Ross for the textbook.